

Importing the necessary libraries

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import warnings, pdftables_api, joblib
warnings.filterwarnings('ignore')
from sklearn.model_selection import train_test_split, cross_val_score, RepeatedKFold, Ra
from sklearn.linear_model import LinearRegression, LassoCV, PassiveAggressiveRegressor,
from sklearn.neighbors import KNeighborsRegressor
from sklearn.svm import SVR, LinearSVR
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor, ExtraTree
from sklearn.neural_network import MLPRegressor
from scipy.stats import probplot
from sklearn.preprocessing import FunctionTransformer, StandardScaler, OrdinalEncoder
from sklearn.impute import KNNImputer
from xgboost import XGBRegressor
from catboost import CatBoostRegressor
from lightgbm import LGBMRegressor
from sklearn.metrics import mean_squared_error, mean_squared_log_error, r2_score, mean_a
from sklearn.feature_selection import SelectKBest, SelectPercentile, SelectFromModel, f_
```

In [2]:

```
## Converting the dataset contained in PDF into a CSV file
# conversion = pdftables_api.Client('zf1L3ni5vlkx')
# conversion.csv('Assignment Dataset - Brandintelle - dataset.pdf', 'dataset.csv')
```

Loading the dataset

In [3]:

```
df = pd.read_csv('dataset.csv', usecols=lambda x: x != 'Sr No')
df.head()
```

Out[3]:

	DATE	Sales	TV_Spends	OOH_Spends	Print_Spends	FB_Impressions	Paid_Search
0	2015-11-23	2754371.667	167687.6	0	95463.66667	7.290385e+07	
1	2015-11-30	2584276.667	214600.9	0	0.00000	1.658110e+07	295
2	2015-12-07	2547386.667	0.0	248022	3404.00000	4.995477e+07	361
3	2015-12-14	2875220.000	625877.3	0	132600.00000	3.164930e+07	368
4	2015-12-21	2215953.333	0.0	520005	0.00000	8.802269e+06	284

Understanding the shape and structure of the dataset

In [4]:

```
df.shape
```

Out[4]:

(208, 11)

In [5]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 208 entries, 0 to 207
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   DATE                  208 non-null   object
1   Sales                 208 non-null   float64
2   TV_Spends             208 non-null   float64
3   OOH_Spends            208 non-null   int64
4   Print_Spends          208 non-null   float64
5   FB_Impressions        208 non-null   float64
6   Paid_Search_Clicks    208 non-null   float64
7   Search_Spends         208 non-null   int64
8   competitor_sales_B    208 non-null   int64
9   FB_Spends             208 non-null   float64
10  events                208 non-null   object
dtypes: float64(6), int64(3), object(2)
memory usage: 18.0+ KB
```

In [6]:

df.describe()

Out[6]:

	Sales	TV_Spends	OOH_Spends	Print_Spends	FB_Impressions	Paid_Sea
count	2.080000e+02	2.080000e+02	208.000000	208.000000	2.080000e+02	208
mean	1.822143e+06	1.113277e+05	81033.639423	27964.741987	2.446024e+07	506
std	7.162286e+05	2.141877e+05	157483.924979	48623.026739	3.509738e+07	406
min	6.722500e+05	0.000000e+00	0.000000	0.000000	0.000000e+00	0
25%	1.165211e+06	0.000000e+00	0.000000	0.000000	0.000000e+00	18
50%	1.874514e+06	0.000000e+00	0.000000	0.000000	0.000000e+00	42
75%	2.378407e+06	1.380503e+05	95359.000000	35758.750002	4.121226e+07	75
max	3.827520e+06	1.185349e+06	938178.000000	239417.333300	1.782983e+08	156

In [7]:

df.isna().sum()

Out[7]:

```

DATE                0
Sales               0
TV_Spends          0
OOH_Spends         0
Print_Spends       0
FB_Impressions     0
Paid_Search_Clicks 0
Search_Spends      0
competitor_sales_B 0
FB_Spends          0
events             0
dtype: int64

```

In [8]:

df.duplicated().sum()

Out[8]:

0

Exploratory Data Analysis


Univariate Analysis

In [9]:

```
for col in df.columns:
    if isinstance(df[col][0], np.number):
        print(f"Skewness of {col}:", df[col].skew())
        print(f"Kurtosis of {col}:", df[col].kurtosis())
        plt.figure(figsize=(14,4))
        plt.subplot(131)
        sns.distplot(df[col])
        plt.subplot(132)
        sns.boxplot(df[col])
        plt.subplot(133)
        probplot(df[col], rvalue=True, plot=plt, dist='norm')
        plt.suptitle(col)
        plt.show()
        plt.tight_layout()
```

Skewness of Sales: 0.24521570453575126

Kurtosis of Sales: -0.8304394212678265



Skewness of TV_Spends: 2.777197561438994

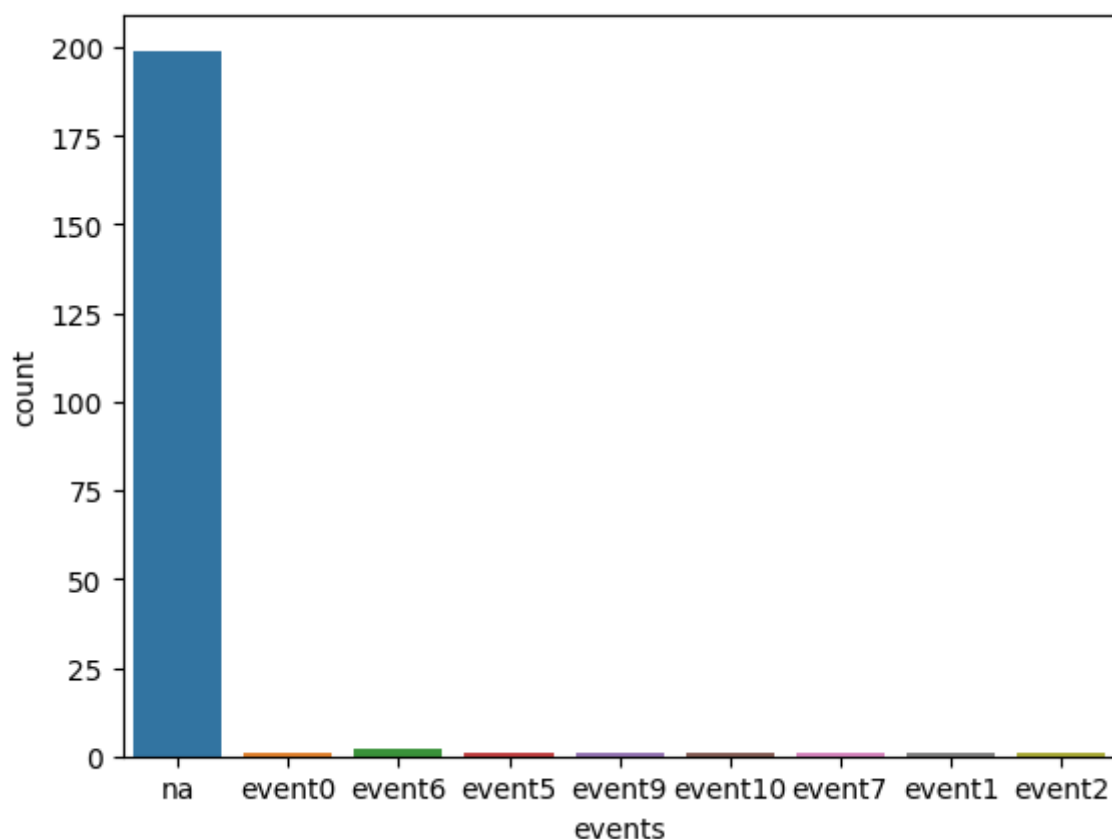
Kurtosis of TV_Spends: 8.654882533891652

<Figure size 640x480 with 0 Axes>

The features "TV_Spends", "OOH_Spends", "Print_Spends", "FB_Impressions" and "FB_Spends" have highly right-skewed distributions due to the presence of outliers.

In [10]:

```
sns.countplot(df.events);
```



Bivariate Analysis

In [11]:

```
df.corr()['Sales'].sort_values(ascending=False)[1:]
```


Out[11]:

competitor_sales_B	0.916454
Search_Spends	0.442774
Paid_Search_Clicks	0.427513
TV_Spends	0.419869
FB_Spends	0.317594
FB_Impressions	0.315126
Print_Spends	0.230415
OOH_Spends	0.095279

Name: Sales, dtype: float64

In [12]:

```
sns.scatterplot(x='competitor_sales_B',y='Sales',data=df);
```



There is a substantially high correlation between competitor sales and the local sales.

In [13]:

```
df.DATE = pd.to_datetime(df.DATE,errors='coerce')
```

In [14]:

```
df['month'] = df.DATE.dt.month_name()  
df['day'] = df.DATE.dt.day_name()  
df['year'] = df.DATE.dt.year  
df['hour'] = df.DATE.dt.hour  
df['daysinmonth'] = df.DATE.dt.daysinmonth  
df['dayofweek'] = df.DATE.dt.dayofweek
```

In [15]:

```
df.describe()
```

Out[15]:

	Sales	TV_Spends	OOH_Spends	Print_Spends	FB_Impressions	Paid_Search
count	2.080000e+02	2.080000e+02	208.000000	208.000000	2.080000e+02	208.000000
mean	1.822143e+06	1.113277e+05	81033.639423	27964.741987	2.446024e+07	50000000.000000
std	7.162286e+05	2.141877e+05	157483.924979	48623.026739	3.509738e+07	40000000.000000
min	6.722500e+05	0.000000e+00	0.000000	0.000000	0.000000e+00	0.000000e+00
25%	1.165211e+06	0.000000e+00	0.000000	0.000000	0.000000e+00	18000000.000000
50%	1.874514e+06	0.000000e+00	0.000000	0.000000	0.000000e+00	42000000.000000
75%	2.378407e+06	1.380503e+05	95359.000000	35758.750002	4.121226e+07	75000000.000000
max	3.827520e+06	1.185349e+06	938178.000000	239417.333300	1.782983e+08	156000000.000000

In [16]:


```
df.drop(['hour', 'dayofweek'], axis=1, inplace=True) # Since the variance of these features
```

In [17]:

```
df['Cost'] = df.TV_Spends + df.OOH_Spends + df.Print_Spends + df.Search_Spends + df.FB_Spends
df['Revenue'] = df.Sales - df.Cost
df.Revenue = df.Revenue.apply(abs)
df['CPM'] = (df.FB_Spends / df.FB_Impressions) * 1000 # Cost Per Thousand Impressions
df['CPC'] = df.Search_Spends / df.Paid_Search_Clicks # Cost Per Click
df['CPS'] = df.Cost / df.Sales # Cost Per Sale
df['ROI'] = (df.Revenue - df.Cost) / (df.Cost+0.0000001) * 100 # Return on Investment
```

In [18]:


```
sns.scatterplot(x='competitor_sales_B',y='Revenue',data=df,color='crimson');
```



There is a significant positive correlation between total revenue and competitor sales.

In [19]:


```
sns.scatterplot(x='CPS',y='OOH_Spends',data=df,color='sienna',s=80,alpha=0.5);
```



There is a strong positive correlation between spends on Outdoor medium and Cost per Sale.

In [20]:


```
sns.scatterplot(x='CPS',y='TV_Spends',data=df,s=100,alpha=0.6,color='darkcyan');
```



There is a strong positive correlation between TV_Spends and Cost Per Sale. Cost per Sale can be minimized by lowering the TV spends.

In [21]:


```
plt.figure(figsize=(10,6))  
sns.barplot(x='events',y='Sales',data=df,palette='coolwarm')  
plt.tight_layout();
```



Event 1 generates the maximum sales whereas event 5 generates the least sales. Moreover, event 0 generates really high sales, less than only event 1. Event 1 and event 0 must be conducted as many times as possible throughout the marketing campaigns to enhance the overall sales.

In [22]:


```
plt.figure(figsize=(12,8))  
sns.barplot(x='month',y='Sales',data=df,palette='ocean');
```



The highest amount of sales are obtained during the final quarter of a year, in the months of October, November and December. A potential reason may be the festive shopping season when people do a lot of shopping due to reasonable discounts and affordable prices. On the other hand, the summer months of April, May and June generate the lowest amount of sales.

In [23]:


```
plt.figure(figsize=(9,6))  
sns.barplot(x='year',y='Sales',data=df,palette='Set2');
```



The year of 2015 recorded the largest sales while the years of 2016 and 2019 observed the lowest sales.

In [24]:


```
plt.figure(figsize=(10,6))  
sns.barplot(x='events',y='Cost',data=df,palette='cubehelix');
```



The highest amount was spent in event 10 whereas the lowest amount was expended in event 7.

In [25]:


```
sns.violinplot(x='year',y='Revenue',data=df,palette='twilight');
```



The year 2015 generated the most revenue whereas 2017 and 2019 were the least profitable years.

In [26]:


```
sns.scatterplot(x='Search_Spends',y='Sales',data=df,color='royalblue',s=100,alpha=0.5);
```



There is a mild positive correlation between spendings done on getting the paid advertisements on the search results and sales i.e. more the amount spent on the paid search results, more will be the total sales.

In [27]:

```
sns.scatterplot(x='Paid_Search_Clicks',y='Sales',data=df,color='tomato',s=100,alpha=0.76
```

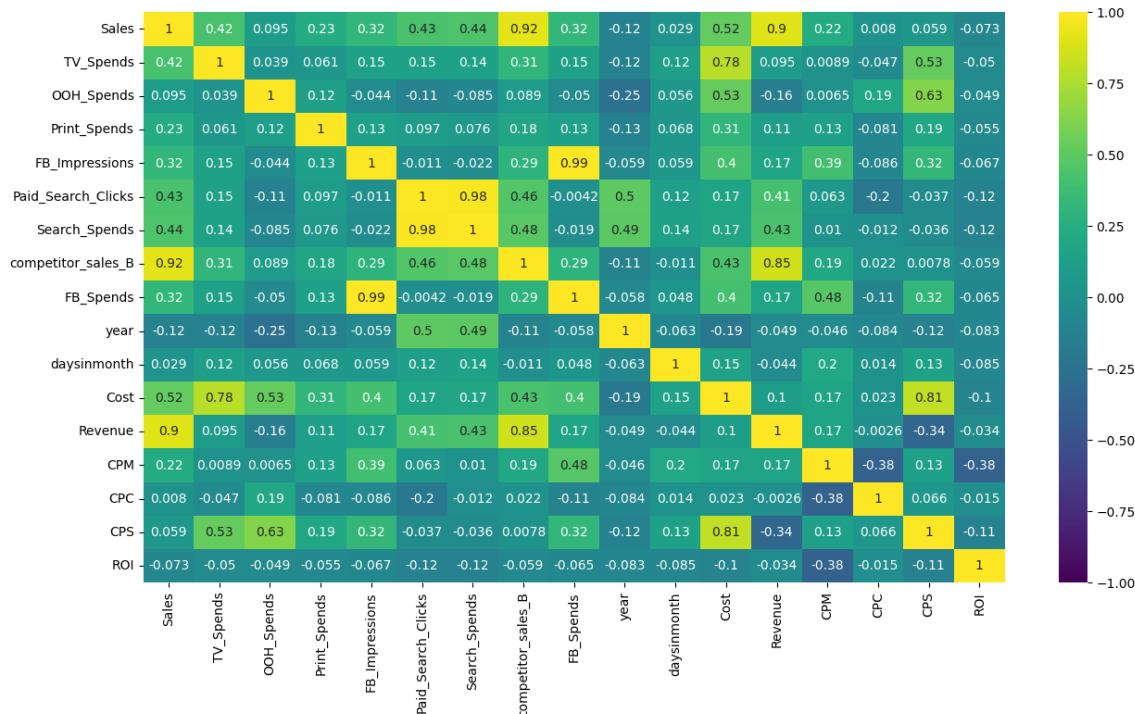


There is a mild positive correlation between sales and the number of user clicks on the paid advertisements posted on the search results i.e. more the user clicks on the paid search results, more will be the gross sales.

Multivariate Analysis

In [28]:

```
plt.figure(figsize=(15,8))
sns.heatmap(df.corr(),annot=True,cmap='viridis',vmin=-1.0,vmax=1.0);
```




There is a perfect positive correlation between the features "FB_impressions" and "FB_Spends" i.e. more the spendings done on the FB advertisements, more will be the total number of user impressions. In addition, there is an almost perfect positive correlation between the features "Paid_Search_Clicks" and "Search_Spends" i.e. more the spendings done on the paid advertisements to place them in the search results, more will be the number of clicks.

There is a mild positive correlation between competitor sales and the spendings done on paid advertisements to get them on the search results as well as number of user clicks on the paid advertisements.

In [29]:

```
sns.clustermap(df.corr());
```



Feature Engineering


Feature Transformation

In [30]:

```
for col in df.columns:
    if isinstance(df[col][0], np.number):
        print(f"Skewness of {col}:", df[col].skew())
        print(f"Kurtosis of {col}:", df[col].kurtosis())
        plt.figure(figsize=(14,4))
        plt.subplot(131)
        sns.distplot(df[col])
        plt.subplot(132)
        sns.boxplot(df[col])
        plt.subplot(133)
        probplot(df[col], rvalue=True, plot=plt, dist='norm')
        plt.suptitle(col)
        plt.show()
        plt.tight_layout()
```

Skewness of Sales: 0.24521570453575126

Kurtosis of Sales: -0.8304394212678265



Skewness of TV_Spends: 2.777197561438994

Kurtosis of TV_Spends: 8.654882533891652

<Figure size 640x480 with 0 Axes>

In [31]:

```
outlier_cols = ['TV_Spends', 'OOH_Spends', 'Print_Spends', 'FB_Impressions', 'Search_Spends']
```

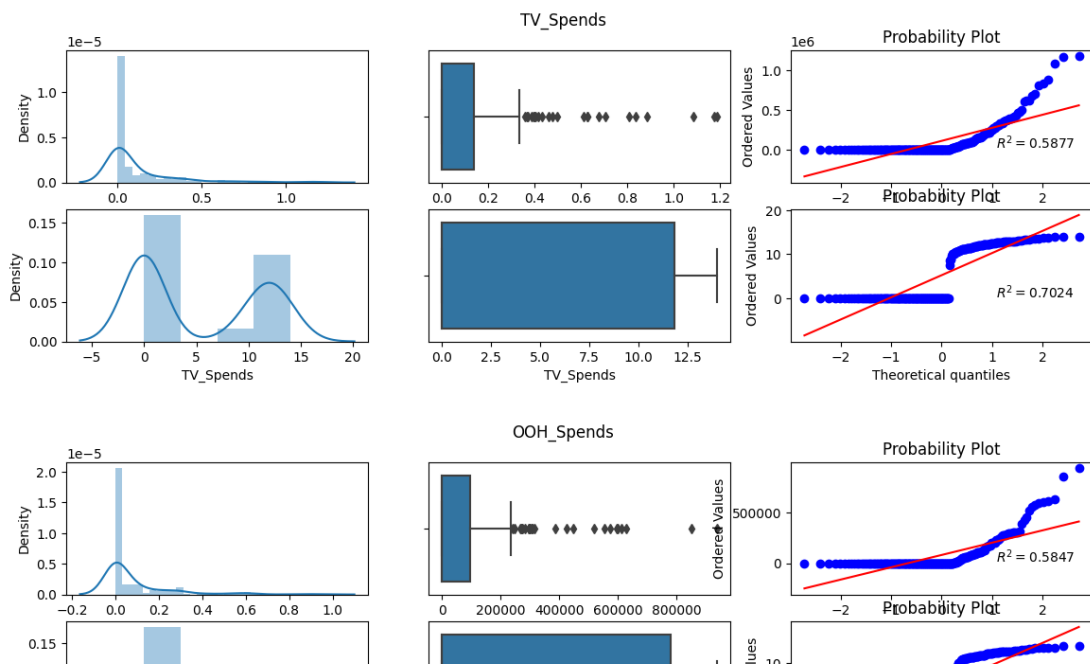
In [32]:

```
def apply_transform(transformer,col):
    plt.figure(figsize=(14,4))
    plt.subplot(231)
    sns.distplot(df[col])
    plt.subplot(232)
    sns.boxplot(df[col])
    plt.subplot(233)
    probplot(df[col],plot=plt,rvalue=True,dist='norm')
    plt.suptitle(col)
    tf_col = transformer.fit_transform(df[col])
    plt.subplot(234)
    sns.distplot(tf_col)
    plt.subplot(235)
    sns.boxplot(tf_col)
    plt.subplot(236)
    probplot(tf_col,plot=plt,rvalue=True,dist='norm')
    plt.show();
```

In [33]:

```
log_transformer = FunctionTransformer(np.log1p)

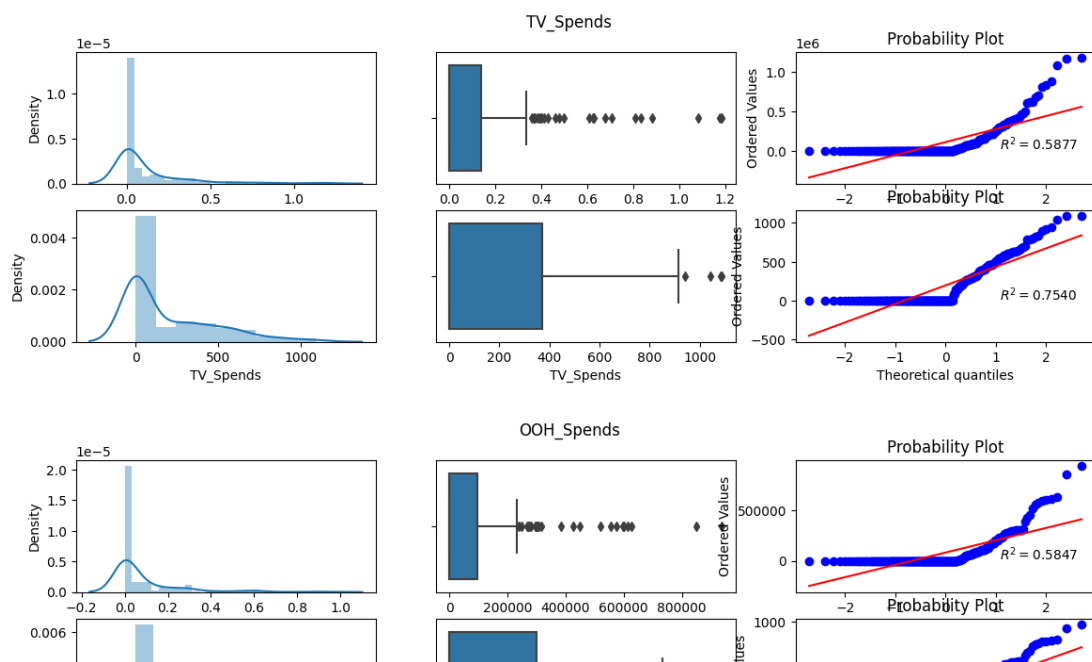
for col in outlier_cols:
    apply_transform(log_transformer,col)
```



In [34]:

```
sqrtransformer = FunctionTransformer(np.sqrt)
```

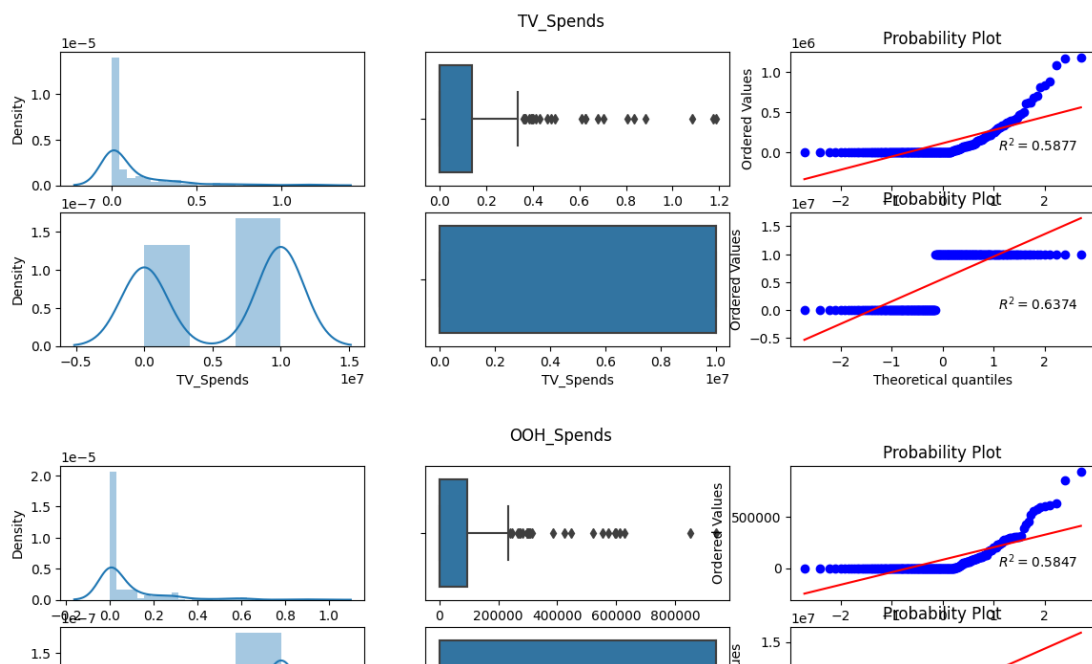
```
for col in outlier_cols:
    apply_transform(sqrtransformer,col)
```



In [35]:

```
reciprocal_transformer = FunctionTransformer(lambda x: 1/(x+0.0000001))
```

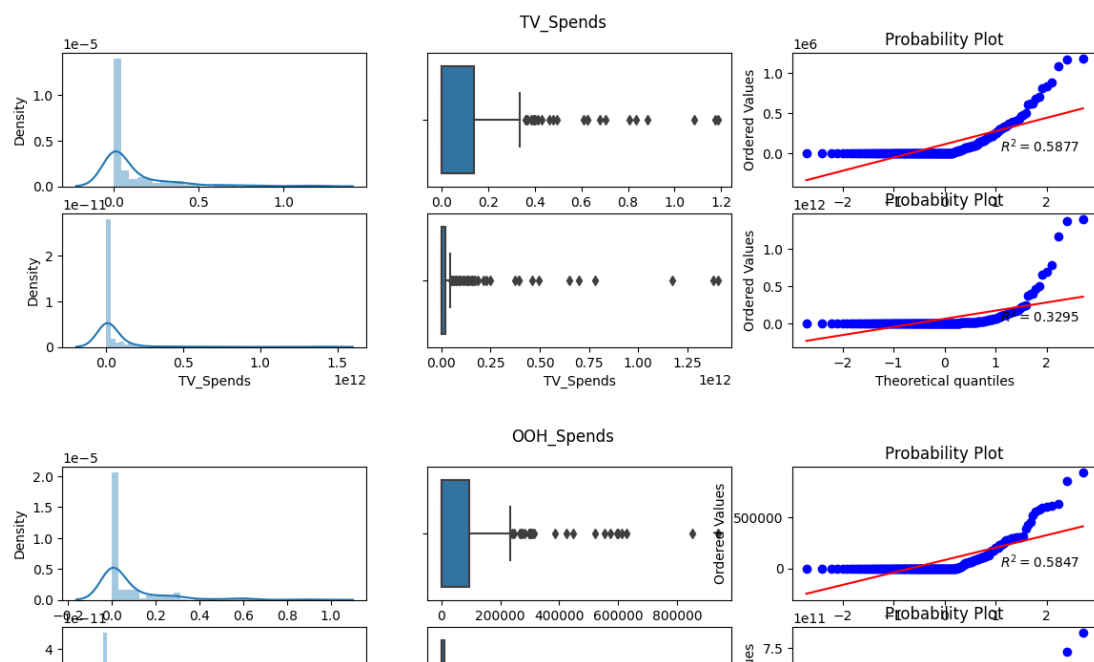
```
for col in outlier_cols:
    apply_transform(reciprocal_transformer,col)
```



In [36]:

```
square_transformer = FunctionTransformer(lambda x: x**2)

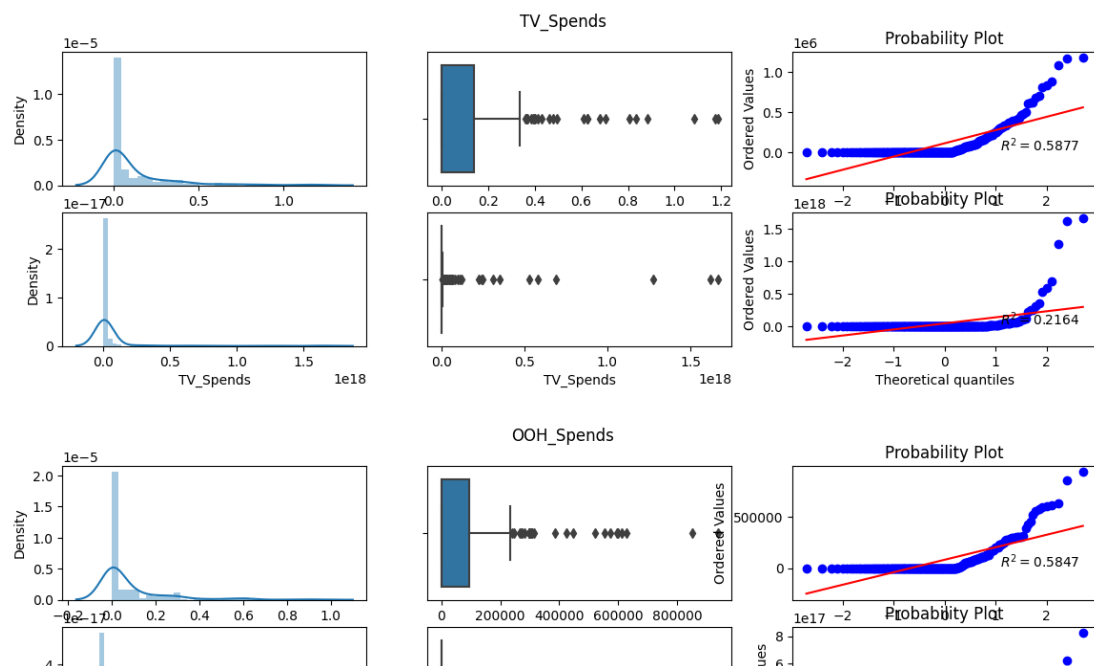
for col in outlier_cols:
    apply_transform(square_transformer,col)
```



In [37]:

```
cube_transformer = FunctionTransformer(lambda x: x**3)

for col in outlier_cols:
    apply_transform(cube_transformer,col)
```



Log Transform: None

Sqrt Transform: TV_Spends, OOH_Spends, Print_Spends, FB_Impressions, Search_Spends, FB_Spends, Cost, CPS

Reciprocal Transform: ROI

In [38]:

```
df.ROI = reciprocal_transformer.fit_transform(df.ROI)
```

In [39]:

```
sqr_cols = ['TV_Spends', 'OOH_Spends', 'Print_Spends', 'FB_Impressions', 'Search_Spends', 'F
for col in sqr_cols:
    df[col] = sqrt_transformer.fit_transform(df[col])
```

In [40]:

```
cube_cols = ['daysinmonth', 'CPM']
for col in cube_cols:
    df[col] = cube_transformer.fit_transform(df[col])
```


Treatment of Outliers

In [41]:

```
for col in df.columns:
    if isinstance(df[col][0], np.number):
        print(f"Skewness of {col}:", df[col].skew())
        print(f"Kurtosis of {col}:", df[col].kurtosis())
        plt.figure(figsize=(14,4))
        plt.subplot(131)
        sns.distplot(df[col])
        plt.subplot(132)
        sns.boxplot(df[col])
        plt.subplot(133)
        probplot(df[col], rvalue=True, plot=plt, dist='norm')
        plt.suptitle(col)
        plt.show()
        plt.tight_layout()
```

Skewness of Sales: 0.24521570453575126

Kurtosis of Sales: -0.8304394212678265



Skewness of TV_Spends: 1.2843203253105715

Kurtosis of TV_Spends: 0.8062730021064808

<Figure size 640x480 with 0 Axes>

In [42]:

```
outlier_cols = ['TV_Spends', 'OOH_Spends', 'Print_Spends', 'daysinmonth', 'Cost', 'CPS', 'ROI']
```

In [43]:

```
def remove_outliers(col):
    lower_limit, upper_limit = df[col].quantile([0.25, 0.75])
    IQR = upper_limit - lower_limit
    lower_whisker = lower_limit - 1.5 * IQR
    upper_whisker = upper_limit + 1.5 * IQR
    return np.where(df[col] > upper_whisker, upper_whisker, np.where(df[col] < lower_whisker, lower_whisker, df[col]))
```

In [44]:


```
for col in outlier_cols:
    df[col] = remove_outliers(col)
```

In [45]:

```
for col in df.columns:
    if isinstance(df[col][0], np.number):
        print(f"Skewness of {col}:", df[col].skew())
        print(f"Kurtosis of {col}:", df[col].kurtosis())
        plt.figure(figsize=(14, 4))
        plt.subplot(131)
        sns.distplot(df[col])
        plt.subplot(132)
        sns.boxplot(df[col])
        plt.subplot(133)
        probplot(df[col], rvalue=True, plot=plt, dist='norm')
        plt.suptitle(col)
        plt.show()
        plt.tight_layout()
```

Skewness of Sales: 0.24521570453575126

Kurtosis of Sales: -0.8304394212678265



Skewness of TV_Spends: 1.1787652374584585

Kurtosis of TV_Spends: 0.2947058456459817

<Figure size 640x480 with 0 Axes>

Categorical Encoding

In [46]:

```
events_encoded = pd.get_dummies(df.events)
df = pd.concat([df,events_encoded],axis=1)
df.head()
```

Out[46]:

	DATE	Sales	TV_Spends	OOH_Spends	Print_Spends	FB_Impressions	Paid_Search
0	2015-11-23	2754371.667	409.496764	0.000000	308.971951	8538.375297	
1	2015-11-30	2584276.667	463.250364	0.000000	0.000000	4071.989630	295
2	2015-12-07	2547386.667	0.000000	498.018072	58.343809	7067.869104	361
3	2015-12-14	2875220.000	791.124074	0.000000	364.142829	5625.770794	368
4	2015-12-21	2215953.333	0.000000	721.113722	0.000000	2966.861891	284

5 rows × 30 columns

In [47]:

```
df.drop(['DATE','events'],axis=1,inplace=True) # Dropping the encoded features
```

In [48]:

```
oe = OrdinalEncoder(categories=[['January','February','March','April','May','June','July']])
df.month = oe.fit_transform(df[['month']])
df.month = df.month.astype(int)
```

In [49]:

```
df.day.unique()
```

Out[49]:

```
array(['Monday'], dtype=object)
```

In [50]:

```
df.drop('day',axis=1,inplace=True)
```

Splitting the dataset into independent and dependent features

In [51]:

```
X = df.drop('Sales',axis=1)
y = df.Sales
```

Dividing the dataset into training and test sets

In [52]:

```
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.3,random_state=101,s
```

Feature Selection

In [53]:

```
# Imputing the missing values of features "CPM" and "CPC" using KNN Imputer
knn = KNNImputer()
X_train.CPM = knn.fit_transform(X_train[['CPM']])
X_test.CPM = knn.transform(X_test[['CPM']])
X_train.CPC = knn.fit_transform(X_train[['CPC']])
X_test.CPC = knn.transform(X_test[['CPC']])
```

In [54]:

```
kbest = SelectKBest(k=10,score_func=f_regression)
kbest.fit(X_train,y_train)
```

Out[54]:

```
SelectKBest(score_func=<function f_regression at 0x000002D852B68C10>)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [55]:

```
selected_features = kbest.get_feature_names_out()
selected_features
```

Out[55]:

```
array(['TV_Spends', 'Print_Spends', 'FB_Impressions',
      'Paid_Search_Clicks', 'Search_Spends', 'competitor_sales_B',
      'FB_Spends', 'month', 'Cost', 'Revenue'], dtype=object)
```

In [56]:

```
p_values = sorted(kbest.pvalues_) [-11:-1]
p_values
```

Out[56]:

```
[0.29621362809634305,
 0.4117374290552849,
 0.4854282592677701,
 0.5191069637621774,
 0.5240722060102341,
 0.5260145210689713,
 0.5610557516486525,
 0.6247269979130835,
 0.7174908169529665,
 0.8149251225529097]
```

In [57]:

```
percentile = SelectPercentile(percentile=40,score_func=mutual_info_regression)
percentile.fit(X_train,y_train)
```

Out[57]:

```
SelectPercentile(percentile=40,
                  score_func=<function mutual_info_regression at 0x000002D8
52B6B640>)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [58]:

```
selected_features = percentile.get_feature_names_out()
selected_features
```

Out[58]:

```
array(['FB_Impressions', 'Paid_Search_Clicks', 'Search_Spends',
      'competitor_sales_B', 'FB_Spends', 'month', 'daysinmonth', 'Cost',
      'Revenue', 'CPC'], dtype=object)
```

In [59]:

```
scores = percentile.scores_
scores
```

Out[59]:

```
array([7.23428038e-02, 0.00000000e+00, 0.00000000e+00, 1.38218084e-01,
      2.26453768e-01, 3.91721631e-01, 1.85563034e+00, 1.40362830e-01,
      1.03035418e+00, 0.00000000e+00, 2.08168839e-01, 1.57324862e-01,
      9.97497534e-01, 7.60956266e-02, 7.85898066e-02, 2.71179576e-03,
      0.00000000e+00, 1.11022302e-16, 2.22044605e-16, 2.29885057e-03,
      2.29885057e-03, 0.00000000e+00, 0.00000000e+00, 6.89655172e-03,
      0.00000000e+00, 1.11484289e-02])
```

In [60]:

```
sfm = SelectFromModel(estimator=RandomForestRegressor(),threshold='median',max_features=sfm.fit(X_train,y_train)
```

Out[60]:

```
SelectFromModel(estimator=RandomForestRegressor(), max_features=10,  
                threshold='median')
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [61]:


```
selected_features = sfm.get_feature_names_out()  
selected_features
```

Out[61]:

```
array(['TV_Spends', 'competitor_sales_B', 'FB_Spends', 'month', 'Cost',  
      'Revenue', 'CPM', 'CPC', 'CPS', 'ROI'], dtype=object)
```

In [62]:

```
rfe = RFECV(estimator=RandomForestRegressor(), step=2, min_features_to_select=10, cv=Repeat  
rfe.fit(X_train, y_train)
```



Fitting estimator with 26 features.
Fitting estimator with 24 features.
Fitting estimator with 22 features.
Fitting estimator with 20 features.
Fitting estimator with 18 features.
Fitting estimator with 16 features.
Fitting estimator with 14 features.
Fitting estimator with 12 features.
Fitting estimator with 26 features.
Fitting estimator with 24 features.
Fitting estimator with 22 features.
Fitting estimator with 20 features.
Fitting estimator with 18 features.
Fitting estimator with 16 features.
Fitting estimator with 14 features.
Fitting estimator with 12 features.
Fitting estimator with 26 features.
Fitting estimator with 24 features.
Fitting estimator with 22 features.
Fitting estimator with 20 features.
Fitting estimator with 18 features.
Fitting estimator with 16 features.
Fitting estimator with 14 features.
Fitting estimator with 12 features.
Fitting estimator with 26 features.
Fitting estimator with 24 features.
Fitting estimator with 22 features.
Fitting estimator with 20 features.
Fitting estimator with 18 features.
Fitting estimator with 16 features.
Fitting estimator with 14 features.
Fitting estimator with 12 features.
Fitting estimator with 26 features.
Fitting estimator with 24 features.
Fitting estimator with 22 features.
Fitting estimator with 20 features.
Fitting estimator with 18 features.
Fitting estimator with 16 features.
Fitting estimator with 14 features.
Fitting estimator with 12 features.
Fitting estimator with 26 features.
Fitting estimator with 24 features.
Fitting estimator with 22 features.
Fitting estimator with 20 features.
Fitting estimator with 18 features.
Fitting estimator with 16 features.
Fitting estimator with 14 features.
Fitting estimator with 12 features.
Fitting estimator with 26 features.
Fitting estimator with 24 features.
Fitting estimator with 22 features.
Fitting estimator with 20 features.
Fitting estimator with 18 features.

```
Fitting estimator with 16 features.  
Fitting estimator with 14 features.  
Fitting estimator with 12 features.  
Fitting estimator with 26 features.  
Fitting estimator with 24 features.  
Fitting estimator with 22 features.  
Fitting estimator with 20 features.  
Fitting estimator with 18 features.  
Fitting estimator with 16 features.  
Fitting estimator with 14 features.  
Fitting estimator with 12 features.  
Fitting estimator with 26 features.  
Fitting estimator with 24 features.  
Fitting estimator with 22 features.  
Fitting estimator with 20 features.  
Fitting estimator with 18 features.  
Fitting estimator with 16 features.  
Fitting estimator with 14 features.  
Fitting estimator with 12 features.  
Fitting estimator with 26 features.  
Fitting estimator with 24 features.  
Fitting estimator with 22 features.  
Fitting estimator with 20 features.  
Fitting estimator with 18 features.  
Fitting estimator with 16 features.  
Fitting estimator with 14 features.  
Fitting estimator with 12 features.
```

Out[62]:

```
RFEVCV(cv=RepeatedKfold(n_repeats=2, n_splits=5, random_state=None),  
       estimator=RandomForestRegressor(), min_features_to_select=10, step=  
2,  
       verbose=1)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [63]:

```
selected_features = rfe.get_feature_names_out()  
selected_features
```

Out[63]:

```
array(['TV_Spends', 'Print_Spends', 'FB_Impressions',  
      'competitor_sales_B', 'month', 'Cost', 'Revenue', 'CPM', 'CPS',  
      'ROI'], dtype=object)
```

In [64]:

```
cv_results = pd.DataFrame(rfe.cv_results_)
cv_results
```

Out[64]:

	mean_test_score	std_test_score	split0_test_score	split1_test_score	split2_test_score	sp
0	0.953846	0.025036	0.934461	0.952460	0.940314	
1	0.952678	0.024481	0.930628	0.952234	0.948455	
2	0.949106	0.026685	0.936129	0.947903	0.933233	
3	0.949527	0.023309	0.927902	0.955995	0.933202	
4	0.950482	0.025952	0.931175	0.948924	0.935897	
5	0.950338	0.027089	0.933802	0.956256	0.925562	
6	0.948114	0.024669	0.938335	0.944549	0.923984	
7	0.947927	0.027466	0.930386	0.943013	0.924109	
8	0.951625	0.027313	0.923922	0.952528	0.942030	

In [65]:

```
sfs = SequentialFeatureSelector(estimator=RandomForestRegressor(),direction='backward',n
sfs.fit(X_train,y_train)
```

Out[65]:

```
SequentialFeatureSelector(direction='backward',
                          estimator=RandomForestRegressor(),
                          n_features_to_select=10)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [66]:

```
selected_features = sfs.get_feature_names_out()
selected_features
```

Out[66]:

```
array(['TV_Spends', 'competitor_sales_B', 'year', 'Cost', 'Revenue',
      'CPS', 'event1', 'event10', 'event5', 'event6'], dtype=object)
```


In [67]:

```
lasso = SelectFromModel(estimator=LassoCV(),max_features=10,threshold='median')
lasso.fit(X_train,y_train)
```

Out[67]:

```
SelectFromModel(estimator=LassoCV(), max_features=10, threshold='median')
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [68]:

```
selected_features = lasso.get_feature_names_out()
selected_features
```

Out[68]:

```
array(['TV_Spends', 'OOH_Spends', 'Print_Spends', 'FB_Impressions',
      'Paid_Search_Clicks', 'Search_Spends', 'competitor_sales_B',
      'FB_Spends', 'month', 'Revenue'], dtype=object)
```

In [69]:

```
rf = RandomForestRegressor()
rf.fit(X_train,y_train)
```

Out[69]:

```
RandomForestRegressor()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [70]:

```

feat_imp = pd.DataFrame({'feature': X_train.columns, 'importance': rf.feature_importance
feat_imp = feat_imp.sort_values('importance',ascending=False)
feat_imp.head(10)

```

Out[70]:

	feature	importance
6	competitor_sales_B	0.526396
12	Revenue	0.421944
11	Cost	0.015886
13	CPM	0.006584
0	TV_Spends	0.006431
8	month	0.004104
14	CPC	0.002922
2	Print_Spends	0.002623
15	CPS	0.002502
7	FB_Spends	0.002249


In [71]:

```

plt.figure(figsize=(10,6))
sns.barplot(x='importance',y='feature',data=feat_imp[:10],orient='horizontal',palette='i
plt.title('Top 10 Features based on Feature Importances',fontsize=32,fontweight='bold',c

```

Top 10 Features based on Feature Importances



Final selected features: TV_Spends, Print_Spends, FB_Impressions, Paid_Search_Clicks, Search_Spends, competitor_sales_B, FB_Spends, month, Cost, Revenue

In [72]:

```
selected_features = kbest.get_feature_names_out()
selected_features
```

Out[72]:

```
array(['TV_Spends', 'Print_Spends', 'FB_Impressions',
      'Paid_Search_Clicks', 'Search_Spends', 'competitor_sales_B',
      'FB_Spends', 'month', 'Cost', 'Revenue'], dtype=object)
```

In [73]:

```
X_train_final = X_train[selected_features]
X_test_final = X_test[selected_features]
```

Feature Scaling and Normalization

In [74]:

```
scaler = StandardScaler()
features = X_train_final.columns
X_train_final = scaler.fit_transform(X_train_final)
X_train_final = pd.DataFrame(X_train_final, columns=features)
X_test_final = scaler.transform(X_test_final)
X_test_final = pd.DataFrame(X_test_final, columns=features)
X_train_final.head()
```

Out[74]:

	TV_Spends	Print_Spends	FB_Impressions	Paid_Search_Clicks	Search_Spends	competitor_sales_B
0	-0.783061	-0.664231	-0.848186	-0.858039	-0.695131	
1	2.711955	1.312909	0.890787	-0.573911	-0.296307	
2	-0.783061	-0.664231	-0.848186	0.618792	0.527713	
3	-0.219207	0.466012	-0.848186	2.490902	1.796331	
4	-0.783061	0.620425	1.103090	0.549877	0.798168	

In [75]:

```
X_test_final.head()
```

Out[75]:

	TV_Spends	Print_Spends	FB_Impressions	Paid_Search_Clicks	Search_Spends	competitor_sales_B
0	-0.783061	0.428265	-0.848186	0.332446	0.310150	
1	1.254396	1.819061	-0.848186	1.097669	0.874541	
2	0.979124	0.564416	-0.848186	-1.296380	-1.903236	
3	-0.475982	-0.664231	-0.848186	1.120051	0.984434	
4	-0.783061	-0.664231	-0.848186	-0.294711	0.094439	

Model Training & Evaluation

In [76]:

```
models = []  
r2_scores = []  
rmse_scores = []  
mape_scores = []  
rmsle_scores = []
```

In [77]:

```
# Custom function for training and evaluating the performance of machine learning algo  
def train_and_evaluate_model(model):  
    models.append(model)  
    model.fit(X_train_final,y_train)  
    pred = model.predict(X_test_final)  
    mape = mean_absolute_percentage_error(y_test,pred)  
    print("Mean Absolute Percentage Error:",mape)  
    mape_scores.append(mape)  
    rmse = np.sqrt(mean_squared_error(y_test,pred))  
    print("Root Mean Squared Error:",rmse)  
    rmse_scores.append(rmse)  
    rmsle = np.sqrt(mean_squared_log_error(y_test,pred))  
    print("Root Mean Squared Log Error:",rmsle)  
    rmsle_scores.append(rmsle)  
    r2 = r2_score(y_test,pred)  
    print("R2 Score:",r2)  
    r2_scores.append(r2)
```

In [78]:

```
train_and_evaluate_model(LinearRegression())
```

Mean Absolute Percentage Error: 0.037602824334626514
Root Mean Squared Error: 95719.65932929565
Root Mean Squared Log Error: 0.054083587741203705
R2 Score: 0.9776935128173896

In [79]:

```
train_and_evaluate_model(RidgeCV())
```

Mean Absolute Percentage Error: 0.037576727144972885
Root Mean Squared Error: 95459.3213232264
Root Mean Squared Log Error: 0.05398820597383878
R2 Score: 0.9778146860267136

In [80]:

```
train_and_evaluate_model(LassoCV())
```

Mean Absolute Percentage Error: 0.03783084232199294
Root Mean Squared Error: 94616.6682867993
Root Mean Squared Log Error: 0.05441759560536869
R2 Score: 0.9782046324533435

In [81]:

```
train_and_evaluate_model(ElasticNetCV())
```

Mean Absolute Percentage Error: 0.387978940632649
Root Mean Squared Error: 637848.9992897833
Root Mean Squared Log Error: 0.3986055426448395
R2 Score: 0.009476800952389897

In [82]:

```
train_and_evaluate_model(PassiveAggressiveRegressor())
```

Mean Absolute Percentage Error: 0.907267207643512
Root Mean Squared Error: 1796564.495941118
Root Mean Squared Log Error: 2.4916878409008674
R2 Score: -6.8580451575580055

In [83]:

```
train_and_evaluate_model(SGDRegressor())
```

Mean Absolute Percentage Error: 0.037566835227451675
Root Mean Squared Error: 94392.6483730361
Root Mean Squared Log Error: 0.0537273258228889
R2 Score: 0.9783077182264387

In [84]:

```
train_and_evaluate_model(HuberRegressor())
```

Mean Absolute Percentage Error: 0.03257423090991409
Root Mean Squared Error: 97713.81345166509
Root Mean Squared Log Error: 0.05412793580298358
R2 Score: 0.9767543968080068

In [85]:

```
train_and_evaluate_model(GammaRegressor())
```

Mean Absolute Percentage Error: 0.12822104797443734
Root Mean Squared Error: 226768.91995867676
Root Mean Squared Log Error: 0.1575623544267843
R2 Score: 0.8748024132026497

In [86]:

```
train_and_evaluate_model(PoissonRegressor())
```

Mean Absolute Percentage Error: 0.06294067286451101
Root Mean Squared Error: 135480.1809989763
Root Mean Squared Log Error: 0.08324090746272196
R2 Score: 0.9553130791185408

In [87]:

```
train_and_evaluate_model(QuantileRegressor())
```

Mean Absolute Percentage Error: 0.39378914922931957
Root Mean Squared Error: 641197.4596635281
Root Mean Squared Log Error: 0.4024142776336646
R2 Score: -0.000950223165545161

In [88]:

```
train_and_evaluate_model(TweedieRegressor())
```

Mean Absolute Percentage Error: 0.12217514530929383
Root Mean Squared Error: 203452.77840820202
Root Mean Squared Log Error: 0.15316495344984618
R2 Score: 0.8992242254554426

In [89]:

```
train_and_evaluate_model(RANSACRegressor())
```

Mean Absolute Percentage Error: 0.03760282433462657
Root Mean Squared Error: 95719.65932929545
Root Mean Squared Log Error: 0.05408358774120371
R2 Score: 0.9776935128173897

In [90]:

```
train_and_evaluate_model(TheilSenRegressor())
```

Mean Absolute Percentage Error: 0.030573411988741174
Root Mean Squared Error: 100104.28013287541
Root Mean Squared Log Error: 0.051124515831031954
R2 Score: 0.9756031257114411

In [91]:

```
train_and_evaluate_model(ARDRegression())
```

Mean Absolute Percentage Error: 0.3897500536465553
Root Mean Squared Error: 640895.4233972202
Root Mean Squared Log Error: 0.40018403325881324
R2 Score: -7.4494259607593705e-06

In [92]:

```
train_and_evaluate_model(KNeighborsRegressor())
```

Mean Absolute Percentage Error: 0.09621601222885663
Root Mean Squared Error: 194296.93162217078
Root Mean Squared Log Error: 0.11857938575928037
R2 Score: 0.9080904205389178

In [93]:

```
train_and_evaluate_model(SVR())
```

Mean Absolute Percentage Error: 0.3937790005193053
Root Mean Squared Error: 641184.4895068462
Root Mean Squared Log Error: 0.40240644405243603
R2 Score: -0.0009097290877075892

In [94]:

```
train_and_evaluate_model(LinearSVR())
```

Mean Absolute Percentage Error: 0.9999072672076434
Root Mean Squared Error: 1932579.9619044955
Root Mean Squared Log Error: 9.369523419245049
R2 Score: -8.092929816658131

In [95]:

```
train_and_evaluate_model(DecisionTreeRegressor())
```

Mean Absolute Percentage Error: 0.05747376450124135
Root Mean Squared Error: 178053.0501159588
Root Mean Squared Log Error: 0.08933740009013505
R2 Score: 0.9228159203428141

In [96]:

```
train_and_evaluate_model(RandomForestRegressor())
```

Mean Absolute Percentage Error: 0.031220698049212087
Root Mean Squared Error: 94152.49581950194
Root Mean Squared Log Error: 0.04413921764540404
R2 Score: 0.9784179562594475

In [97]:

```
train_and_evaluate_model(BaggingRegressor())
```

Mean Absolute Percentage Error: 0.03229415620859832
Root Mean Squared Error: 93314.61261437052
Root Mean Squared Log Error: 0.0459340413587151
R2 Score: 0.9788003734975818

In [98]:

```
train_and_evaluate_model(ExtraTreesRegressor())
```

Mean Absolute Percentage Error: 0.020760607330991044
Root Mean Squared Error: 80884.65690320772
Root Mean Squared Log Error: 0.04130882721501735
R2 Score: 0.9840720020957423

In [99]:

```
train_and_evaluate_model(GradientBoostingRegressor())
```

Mean Absolute Percentage Error: 0.025872045660698333
Root Mean Squared Error: 79388.0408617185
Root Mean Squared Log Error: 0.04022760523143929
R2 Score: 0.9846559832574224

In [100]:

```
train_and_evaluate_model(AdaBoostRegressor())
```

Mean Absolute Percentage Error: 0.06038815051975686
Root Mean Squared Error: 130780.25885881713
Root Mean Squared Log Error: 0.07690470681856508
R2 Score: 0.958359754639006

In [101]:

```
train_and_evaluate_model(HistGradientBoostingRegressor())
```

Mean Absolute Percentage Error: 0.04854479957493201
Root Mean Squared Error: 133895.73833799115
Root Mean Squared Log Error: 0.0674887118995013
R2 Score: 0.9563521955001669

In [102]:

```
train_and_evaluate_model(MLPRegressor())
```

Mean Absolute Percentage Error: 0.9999917598610703
Root Mean Squared Error: 1932698.5618114949
Root Mean Squared Log Error: 11.735067605729451
R2 Score: -8.09404589334577

In [103]:

```
train_and_evaluate_model(XGBRegressor())
```

Mean Absolute Percentage Error: 0.042062524512644826
Root Mean Squared Error: 128363.67129137942
Root Mean Squared Log Error: 0.06035237338180336
R2 Score: 0.959884412763658

In [104]:

```
train_and_evaluate_model(CatBoostRegressor(silent=True))
```

Mean Absolute Percentage Error: 0.03515162308993974
Root Mean Squared Error: 108755.73903371001
Root Mean Squared Log Error: 0.061378420723152506
R2 Score: 0.9712039272265389

In [105]:

```
train_and_evaluate_model(LGBMRegressor())
```

Mean Absolute Percentage Error: 0.04485766673147905
Root Mean Squared Error: 126248.80450768779
Root Mean Squared Log Error: 0.06491580174683761
R2 Score: 0.9611953792536151

In [106]:

```
train_and_evaluate_model(VotingRegressor(estimators=[  
    ('CAT', CatBoostRegressor(silent=True)),  
    ('LR', LinearRegression()),  
    ('SGD', SGDRegressor()),  
    ('GB', GradientBoostingRegressor()),  
    ('ET', ExtraTreesRegressor()),  
    ('BAG', BaggingRegressor()),  
    ('RF', RandomForestRegressor()),  
    ('THEIL', TheilSenRegressor()),  
    ('RANSAC', RANSACRegressor()),  
    ('HUBER', HuberRegressor()),  
    ('LASSO', LassoCV()),  
    ('RIDGE', RidgeCV())  
]))
```

Mean Absolute Percentage Error: 0.02410496728707758
Root Mean Squared Error: 74835.45294225398
Root Mean Squared Log Error: 0.03774425811193272
R2 Score: 0.9863653599329305

Baseline Model Performance Comparison

In [107]:

```
model_perfs = pd.DataFrame({'Model': models, 'MAPE': mape_scores, 'RMSE': rmse_scores, 'R2': r2_scores})
model_perfs = model_perfs.sort_values('R2', ascending=False)
model_perfs
```

Out[107]:

	Model	MAPE	RMSE	RMSLE	
28	VotingRegressor(estimators=[('CAT', ln ...	0.024105	7.483545e+04	0.037744	0.98
21	([DecisionTreeRegressor(criterion='friedman_ms...	0.025872	7.938804e+04	0.040228	0.98
20	(ExtraTreeRegressor(random_state=2090432581), ...	0.020761	8.088466e+04	0.041309	0.98
19	(DecisionTreeRegressor(random_state=174724458)...	0.032294	9.331461e+04	0.045934	0.97
18	(DecisionTreeRegressor(max_features=1.0, rando...	0.031221	9.415250e+04	0.044139	0.97
5	SGDRegressor()	0.037567	9.439265e+04	0.053727	0.97
2	LassoCV()	0.037831	9.461667e+04	0.054418	0.97
1	RidgeCV()	0.037577	9.545932e+04	0.053988	0.97
11	RANSACRegressor()	0.037603	9.571966e+04	0.054084	0.97
0	LinearRegression()	0.037603	9.571966e+04	0.054084	0.97
6	HuberRegressor()	0.032574	9.771381e+04	0.054128	0.97
12	TheilSenRegressor()	0.030573	1.001043e+05	0.051125	0.97
26	<catboost.core.CatBoostRegressor object at 0x0...	0.035152	1.087557e+05	0.061378	0.97
27	LGBMRegressor()	0.044858	1.262488e+05	0.064916	0.96
25	XGBRegressor(base_score=None, booster=None, ca...	0.042063	1.283637e+05	0.060352	0.95
22	(DecisionTreeRegressor(max_depth=3, random_sta...	0.060388	1.307803e+05	0.076905	0.95
23	HistGradientBoostingRegressor()	0.048545	1.338957e+05	0.067489	0.95
8	PoissonRegressor()	0.062941	1.354802e+05	0.083241	0.95
17	DecisionTreeRegressor()	0.057474	1.780531e+05	0.089337	0.92
14	KNeighborsRegressor()	0.096216	1.942969e+05	0.118579	0.90
10	TweedieRegressor()	0.122175	2.034528e+05	0.153165	0.89
7	GammaRegressor()	0.128221	2.267689e+05	0.157562	0.87
3	ElasticNetCV()	0.387979	6.378490e+05	0.398606	0.00
13	ARDRegression()	0.389750	6.408954e+05	0.400184	-0.00
15	SVR()	0.393779	6.411845e+05	0.402406	-0.00
9	QuantileRegressor()	0.393789	6.411975e+05	0.402414	-0.00
4	PassiveAggressiveRegressor()	0.907267	1.796564e+06	2.491688	-6.85
16	LinearSVR()	0.999907	1.932580e+06	9.369523	-8.09
24	MLPRegressor()	0.999992	1.932699e+06	11.735068	-8.09

Among the baseline models, the Voting Regressor produced the best performance by achieving the maximum r^2 score of more than 98.5%, closely followed by Extra Trees Regressor and Decision Tree Regressor, which obtained r^2 scores of 98.47% and 98.32% respectively.

Cross Validation and Model Optimization using Hyperparameter Tuning

In [108]:

```
param_grid = {'n_estimators': [200,400,600,800,1000],
              'criterion': ['squared_error', 'absolute_error', 'friedman_mse', 'poisson'],
              'max_features': ['auto','sqrt','log2'],
              'bootstrap': [True,False],
              'oob_score': [True,False]}

grid_rf = RandomizedSearchCV(RandomForestRegressor(),param_grid,verbose=2,cv=5)
train_and_evaluate_model(grid_rf)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[CV] END bootstrap=False, criterion=squared_error, max_features=sqrt, n_estimators=400, oob_score=False; total time= 1.1s
[CV] END bootstrap=False, criterion=squared_error, max_features=sqrt, n_estimators=400, oob_score=False; total time= 1.1s
[CV] END bootstrap=False, criterion=squared_error, max_features=sqrt, n_estimators=400, oob_score=False; total time= 1.1s
[CV] END bootstrap=False, criterion=squared_error, max_features=sqrt, n_estimators=400, oob_score=False; total time= 1.1s
[CV] END bootstrap=False, criterion=squared_error, max_features=sqrt, n_estimators=400, oob_score=False; total time= 1.1s
[CV] END bootstrap=True, criterion=absolute_error, max_features=auto, n_estimators=600, oob_score=True; total time= 5.0s
[CV] END bootstrap=True, criterion=absolute_error, max_features=auto, n_estimators=600, oob_score=True; total time= 5.2s
[CV] END bootstrap=True, criterion=absolute_error, max_features=auto, n_estimators=600, oob_score=True; total time= 5.0s
[CV] END bootstrap=True, criterion=absolute_error, max_features=auto, n_estimators=600, oob_score=True; total time= 4.9s
[CV] END bootstrap=True, criterion=absolute_error, max_features=auto, n_estimators=600, oob_score=True; total time= 5.0s
[CV] END bootstrap=False, criterion=squared_error, max_features=log2, n_estimators=200, oob_score=False; total time= 0.5s
[CV] END bootstrap=False, criterion=squared_error, max_features=log2, n_estimators=200, oob_score=False; total time= 0.5s
[CV] END bootstrap=False, criterion=squared_error, max_features=log2, n_estimators=200, oob_score=False; total time= 0.5s
[CV] END bootstrap=False, criterion=squared_error, max_features=log2, n_estimators=200, oob_score=False; total time= 0.5s
[CV] END bootstrap=False, criterion=squared_error, max_features=log2, n_estimators=200, oob_score=False; total time= 0.5s
[CV] END bootstrap=False, criterion=squared_error, max_features=auto, n_estimators=1000, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, criterion=squared_error, max_features=auto, n_estimators=1000, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, criterion=squared_error, max_features=auto, n_estimators=1000, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, criterion=squared_error, max_features=auto, n_estimators=1000, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, criterion=squared_error, max_features=auto, n_estimators=1000, oob_score=True; total time= 0.0s
[CV] END bootstrap=True, criterion=absolute_error, max_features=log2, n_estimators=800, oob_score=True; total time= 4.5s
[CV] END bootstrap=True, criterion=absolute_error, max_features=log2, n_estimators=800, oob_score=True; total time= 4.8s
[CV] END bootstrap=True, criterion=absolute_error, max_features=log2, n_estimators=800, oob_score=True; total time= 5.1s
[CV] END bootstrap=True, criterion=absolute_error, max_features=log2, n_estimators=800, oob_score=True; total time= 4.4s
[CV] END bootstrap=True, criterion=absolute_error, max_features=log2, n_estimators=800, oob_score=True; total time= 4.6s
[CV] END bootstrap=False, criterion=poisson, max_features=sqrt, n_estimators=400, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, criterion=poisson, max_features=sqrt, n_estimators=400, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, criterion=poisson, max_features=sqrt, n_estimators=400, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, criterion=poisson, max_features=sqrt, n_estimators=400, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, criterion=poisson, max_features=sqrt, n_estimators=400, oob_score=True; total time= 0.0s
```

```
[CV] END bootstrap=True, criterion=friedman_mse, max_features=log2, n_estimators=200, oob_score=False; total time= 0.6s
[CV] END bootstrap=True, criterion=friedman_mse, max_features=log2, n_estimators=200, oob_score=False; total time= 0.7s
[CV] END bootstrap=True, criterion=friedman_mse, max_features=log2, n_estimators=200, oob_score=False; total time= 0.6s
[CV] END bootstrap=True, criterion=friedman_mse, max_features=log2, n_estimators=200, oob_score=False; total time= 0.6s
[CV] END bootstrap=True, criterion=friedman_mse, max_features=log2, n_estimators=200, oob_score=False; total time= 0.7s
[CV] END bootstrap=True, criterion=absolute_error, max_features=log2, n_estimators=200, oob_score=False; total time= 1.0s
[CV] END bootstrap=True, criterion=absolute_error, max_features=log2, n_estimators=200, oob_score=False; total time= 1.0s
[CV] END bootstrap=True, criterion=absolute_error, max_features=log2, n_estimators=200, oob_score=False; total time= 1.0s
[CV] END bootstrap=True, criterion=absolute_error, max_features=log2, n_estimators=200, oob_score=False; total time= 0.9s
[CV] END bootstrap=True, criterion=absolute_error, max_features=log2, n_estimators=200, oob_score=False; total time= 0.9s
[CV] END bootstrap=True, criterion=absolute_error, max_features=sqrt, n_estimators=600, oob_score=False; total time= 3.0s
[CV] END bootstrap=True, criterion=absolute_error, max_features=sqrt, n_estimators=600, oob_score=False; total time= 3.0s
[CV] END bootstrap=True, criterion=absolute_error, max_features=sqrt, n_estimators=600, oob_score=False; total time= 2.9s
[CV] END bootstrap=True, criterion=absolute_error, max_features=sqrt, n_estimators=600, oob_score=False; total time= 3.0s
[CV] END bootstrap=True, criterion=absolute_error, max_features=sqrt, n_estimators=600, oob_score=False; total time= 3.0s
[CV] END bootstrap=False, criterion=friedman_mse, max_features=auto, n_estimators=600, oob_score=False; total time= 2.1s
[CV] END bootstrap=False, criterion=friedman_mse, max_features=auto, n_estimators=600, oob_score=False; total time= 2.1s
[CV] END bootstrap=False, criterion=friedman_mse, max_features=auto, n_estimators=600, oob_score=False; total time= 2.1s
[CV] END bootstrap=False, criterion=friedman_mse, max_features=auto, n_estimators=600, oob_score=False; total time= 2.1s
[CV] END bootstrap=False, criterion=friedman_mse, max_features=auto, n_estimators=600, oob_score=False; total time= 2.2s
Mean Absolute Percentage Error: 0.03207742438582975
Root Mean Squared Error: 96287.085641307
Root Mean Squared Log Error: 0.04740284868926336
R2 Score: 0.9774282631485722
```

In [109]:

```
grid_rf.best_params_
```

Out[109]:

```
{'oob_score': True,
 'n_estimators': 600,
 'max_features': 'auto',
 'criterion': 'absolute_error',
 'bootstrap': True}
```


In [112]:

```
param_grid = {'learning_rate': [0.2,0.4,0.5,0.8,1.0],  
              'loss': ['squared_error', 'absolute_error', 'poisson', 'quantile'],  
              'max_bins': np.arange(0,255,50),  
              'interaction_cst': ['pairwise', 'no_interaction']}  
grid_hgb = RandomizedSearchCV(HistGradientBoostingRegressor(), param_grid, cv=5, verbose=2)  
train_and_evaluate_model(grid_hgb)
```



```
Fitting 5 folds for each of 10 candidates, totalling 50 fits
[CV] END interaction_cst=no_interaction, learning_rate=1.0, loss=quantile,
max_bins=100; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=1.0, loss=quantile,
max_bins=100; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=1.0, loss=quantile,
max_bins=100; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=1.0, loss=quantile,
max_bins=100; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=1.0, loss=quantile,
max_bins=100; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=0.5, loss=absolute_
error, max_bins=150; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=0.5, loss=absolute_
error, max_bins=150; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=0.5, loss=absolute_
error, max_bins=150; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=0.5, loss=absolute_
error, max_bins=150; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=0.5, loss=absolute_
error, max_bins=150; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=0.5, loss=absolute_
error, max_bins=150; total time= 0.0s
[CV] END interaction_cst=pairwise, learning_rate=0.2, loss=absolute_error,
max_bins=100; total time= 0.2s
[CV] END interaction_cst=pairwise, learning_rate=0.2, loss=absolute_error,
max_bins=100; total time= 0.2s
[CV] END interaction_cst=pairwise, learning_rate=0.2, loss=absolute_error,
max_bins=100; total time= 0.2s
[CV] END interaction_cst=pairwise, learning_rate=0.2, loss=absolute_error,
max_bins=100; total time= 0.2s
[CV] END interaction_cst=pairwise, learning_rate=0.2, loss=absolute_error,
max_bins=100; total time= 0.2s
[CV] END interaction_cst=pairwise, learning_rate=0.8, loss=poisson, max_bi
ns=250; total time= 0.1s
[CV] END interaction_cst=pairwise, learning_rate=0.8, loss=poisson, max_bi
ns=250; total time= 0.1s
[CV] END interaction_cst=pairwise, learning_rate=0.8, loss=poisson, max_bi
ns=250; total time= 0.1s
[CV] END interaction_cst=pairwise, learning_rate=0.8, loss=poisson, max_bi
ns=250; total time= 0.1s
[CV] END interaction_cst=pairwise, learning_rate=0.8, loss=poisson, max_bi
ns=250; total time= 0.1s
[CV] END interaction_cst=pairwise, learning_rate=0.2, loss=absolute_error,
max_bins=0; total time= 0.0s
[CV] END interaction_cst=pairwise, learning_rate=0.2, loss=absolute_error,
max_bins=0; total time= 0.0s
[CV] END interaction_cst=pairwise, learning_rate=0.2, loss=absolute_error,
max_bins=0; total time= 0.0s
[CV] END interaction_cst=pairwise, learning_rate=0.2, loss=absolute_error,
max_bins=0; total time= 0.0s
[CV] END interaction_cst=pairwise, learning_rate=0.2, loss=poisson, max_bi
ns=0; total time= 0.0s
[CV] END interaction_cst=pairwise, learning_rate=0.2, loss=poisson, max_bi
ns=0; total time= 0.0s
[CV] END interaction_cst=pairwise, learning_rate=0.2, loss=poisson, max_bi
ns=0; total time= 0.0s
[CV] END interaction_cst=pairwise, learning_rate=0.2, loss=poisson, max_bi
ns=0; total time= 0.0s
[CV] END interaction_cst=pairwise, learning_rate=0.2, loss=poisson, max_bi
ns=0; total time= 0.0s
```

```
[CV] END interaction_cst=no_interaction, learning_rate=0.8, loss=poisson,
max_bins=250; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=0.8, loss=poisson,
max_bins=250; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=0.8, loss=poisson,
max_bins=250; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=0.8, loss=poisson,
max_bins=250; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=0.8, loss=poisson,
max_bins=250; total time= 0.0s
[CV] END interaction_cst=pairwise, learning_rate=0.8, loss=squared_error,
max_bins=0; total time= 0.0s
[CV] END interaction_cst=pairwise, learning_rate=0.8, loss=squared_error,
max_bins=0; total time= 0.0s
[CV] END interaction_cst=pairwise, learning_rate=0.8, loss=squared_error,
max_bins=0; total time= 0.0s
[CV] END interaction_cst=pairwise, learning_rate=0.8, loss=squared_error,
max_bins=0; total time= 0.0s
[CV] END interaction_cst=pairwise, learning_rate=0.8, loss=squared_error,
max_bins=0; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=0.2, loss=absolute_
error, max_bins=250; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=0.2, loss=absolute_
error, max_bins=250; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=0.2, loss=absolute_
error, max_bins=250; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=0.2, loss=absolute_
error, max_bins=250; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=0.2, loss=absolute_
error, max_bins=250; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=0.2, loss=absolute_
error, max_bins=50; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=0.2, loss=absolute_
error, max_bins=50; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=0.2, loss=absolute_
error, max_bins=50; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=0.2, loss=absolute_
error, max_bins=50; total time= 0.0s
[CV] END interaction_cst=no_interaction, learning_rate=0.2, loss=absolute_
error, max_bins=50; total time= 0.0s
Mean Absolute Percentage Error: 0.05068667357430494
Root Mean Squared Error: 137696.6523488204
Root Mean Squared Log Error: 0.07847529290656403
R2 Score: 0.9538389521789743
```

In [113]:

```
grid_hgb.best_params_
```

Out[113]:

```
{'max_bins': 100,
 'loss': 'absolute_error',
 'learning_rate': 0.2,
 'interaction_cst': 'pairwise'}
```

In [114]:

```

param_grid = {'C': [0.001,0.01,0.1,1,5],
              'gamma': ['scale','auto'],
              'epsilon': np.linspace(0.001,1,5),
              'kernel': ['linear','poly','rbf','sigmoid'],
              'degree': [2,3,4,5],
              'shrinking': [True,False]
              }
grid_svr = RandomizedSearchCV(SVR(),param_grid,verbose=2,cv=RepeatedKFold(n_splits=5,n_r
train_and_evaluate_model(grid_svr)

```

Fitting 15 folds for each of 10 candidates, totalling 150 fits

```

[CV] END C=5, degree=3, epsilon=1.0, gamma=scale, kernel=linear, shrink
ing=True; total time= 0.0s
[CV] END C=5, degree=3, epsilon=1.0, gamma=scale, kernel=linear, shrink
ing=True; total time= 0.0s
[CV] END C=5, degree=3, epsilon=1.0, gamma=scale, kernel=linear, shrink
ing=True; total time= 0.0s
[CV] END C=5, degree=3, epsilon=1.0, gamma=scale, kernel=linear, shrink
ing=True; total time= 0.0s
[CV] END C=5, degree=3, epsilon=1.0, gamma=scale, kernel=linear, shrink
ing=True; total time= 0.0s
[CV] END C=5, degree=3, epsilon=1.0, gamma=scale, kernel=linear, shrink
ing=True; total time= 0.0s
[CV] END C=5, degree=3, epsilon=1.0, gamma=scale, kernel=linear, shrink
ing=True; total time= 0.0s
[CV] END C=5, degree=3, epsilon=1.0, gamma=scale, kernel=linear, shrink
ing=True; total time= 0.0s
[CV] END C=5, degree=3, epsilon=1.0, gamma=scale, kernel=linear, shrink
ing=True; total time= 0.0s
[CV] END C=5, degree=3, epsilon=1.0, gamma=scale, kernel=linear, shrink
ing=True; total time= 0.0s

```

In [115]:

```
grid_svr.best_params_
```

Out[115]:

```

{'shrinking': True,
 'kernel': 'linear',
 'gamma': 'scale',
 'epsilon': 1.0,
 'degree': 3,
 'C': 5}

```

In [116]:

```
param_grid = {'loss': ['epsilon_insensitive', 'squared_epsilon_insensitive'],
              'C': [0.0001, 0.001, 0.01, 0.1, 1],
              'epsilon': np.linspace(0.001, 1, 5)}

grid_lsvr = RandomizedSearchCV(LinearSVR(), param_grid, cv=RepeatedKFold(n_splits=6, n_repe
train_and_evaluate_model(grid_lsvr)
```

Fitting 12 folds for each of 10 candidates, totalling 120 fits

```
[CV] END C=0.01, epsilon=0.75025, loss=squared_epsilon_insensitive; tot
al time= 0.0s
[CV] END C=0.01, epsilon=0.75025, loss=squared_epsilon_insensitive; tot
al time= 0.0s
[CV] END C=0.01, epsilon=0.75025, loss=squared_epsilon_insensitive; tot
al time= 0.0s
[CV] END C=0.01, epsilon=0.75025, loss=squared_epsilon_insensitive; tot
al time= 0.0s
[CV] END C=0.01, epsilon=0.75025, loss=squared_epsilon_insensitive; tot
al time= 0.0s
[CV] END C=0.01, epsilon=0.75025, loss=squared_epsilon_insensitive; tot
al time= 0.0s
[CV] END C=0.01, epsilon=0.75025, loss=squared_epsilon_insensitive; tot
al time= 0.0s
[CV] END C=0.01, epsilon=0.75025, loss=squared_epsilon_insensitive; tot
al time= 0.0s
[CV] END C=0.01, epsilon=0.75025, loss=squared_epsilon_insensitive; tot
al time= 0.0s
[CV] END C=0.01, epsilon=0.75025, loss=squared_epsilon_insensitive; tot
al time= 0.0s
```

In [117]:

```
grid_lsvr.best_params_
```

Out[117]:

```
{'loss': 'squared_epsilon_insensitive', 'epsilon': 0.25075, 'C': 0.1}
```


In [120]:

```
param_grid = {'loss': ['squared_error', 'absolute_error', 'huber', 'quantile'],
              'n_estimators': [100, 400, 700, 1000],
              'learning_rate': [0.2, 0.4, 0.7, 1],
              'criterion': ['friedman_mse', 'squared_error'],
              'max_features': ['auto', 'sqrt', 'log2']}

grid_gb = RandomizedSearchCV(GradientBoostingRegressor(), param_grid, verbose=3, cv=6)
train_and_evaluate_model(grid_gb)
```

Fitting 6 folds for each of 10 candidates, totalling 60 fits

[CV 1/6] END criterion=friedman_mse, learning_rate=0.7, loss=squared_error, max_features=sqrt, n_estimators=100;; score=0.941 total time= 0.1s

[CV 2/6] END criterion=friedman_mse, learning_rate=0.7, loss=squared_error, max_features=sqrt, n_estimators=100;; score=0.934 total time= 0.1s

[CV 3/6] END criterion=friedman_mse, learning_rate=0.7, loss=squared_error, max_features=sqrt, n_estimators=100;; score=0.834 total time= 0.0s

[CV 4/6] END criterion=friedman_mse, learning_rate=0.7, loss=squared_error, max_features=sqrt, n_estimators=100;; score=0.691 total time= 0.0s

[CV 5/6] END criterion=friedman_mse, learning_rate=0.7, loss=squared_error, max_features=sqrt, n_estimators=100;; score=0.872 total time= 0.0s

[CV 6/6] END criterion=friedman_mse, learning_rate=0.7, loss=squared_error, max_features=sqrt, n_estimators=100;; score=0.908 total time= 0.1s

In [121]:

```
grid_gb.best_params_
```

Out[121]:

```
{'n_estimators': 100,
 'max_features': 'auto',
 'loss': 'absolute_error',
 'learning_rate': 0.4,
 'criterion': 'squared_error'}
```

In [122]:

```
param_grid = {'learning_rate': [0.2,0.4,0.7,1],  
              'n_estimators': [100,400,700,1000]  
              }  
  
grid_cat = RandomizedSearchCV(CatBoostRegressor(silent=True),param_grid,verbose=5,cv=5)  
train_and_evaluate_model(grid_cat)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[CV 1/5] END learning_rate=0.4, n_estimators=700;; score=0.927 total time=
2.2s
[CV 2/5] END learning_rate=0.4, n_estimators=700;; score=0.961 total time=
2.1s
[CV 3/5] END learning_rate=0.4, n_estimators=700;; score=0.884 total time=
2.0s
[CV 4/5] END learning_rate=0.4, n_estimators=700;; score=0.975 total time=
2.3s
[CV 5/5] END learning_rate=0.4, n_estimators=700;; score=0.941 total time=
3.0s
[CV 1/5] END .learning_rate=1, n_estimators=100;; score=0.887 total time=
1.1s
[CV 2/5] END .learning_rate=1, n_estimators=100;; score=0.919 total time=
1.2s
[CV 3/5] END .learning_rate=1, n_estimators=100;; score=0.815 total time=
1.0s
[CV 4/5] END .learning_rate=1, n_estimators=100;; score=0.935 total time=
0.9s
[CV 5/5] END .learning_rate=1, n_estimators=100;; score=0.920 total time=
1.0s
[CV 1/5] END learning_rate=0.4, n_estimators=400;; score=0.927 total time=
1.7s
[CV 2/5] END learning_rate=0.4, n_estimators=400;; score=0.961 total time=
1.6s
[CV 3/5] END learning_rate=0.4, n_estimators=400;; score=0.884 total time=
1.4s
[CV 4/5] END learning_rate=0.4, n_estimators=400;; score=0.975 total time=
1.7s
[CV 5/5] END learning_rate=0.4, n_estimators=400;; score=0.941 total time=
1.5s
[CV 1/5] END learning_rate=0.2, n_estimators=1000;; score=0.904 total time
= 3.4s
[CV 2/5] END learning_rate=0.2, n_estimators=1000;; score=0.967 total time
= 2.8s
[CV 3/5] END learning_rate=0.2, n_estimators=1000;; score=0.930 total time
= 2.6s
[CV 4/5] END learning_rate=0.2, n_estimators=1000;; score=0.973 total time
= 3.4s
[CV 5/5] END learning_rate=0.2, n_estimators=1000;; score=0.970 total time
= 2.8s
[CV 1/5] END learning_rate=0.2, n_estimators=100;; score=0.902 total time=
1.1s
[CV 2/5] END learning_rate=0.2, n_estimators=100;; score=0.967 total time=
0.9s
[CV 3/5] END learning_rate=0.2, n_estimators=100;; score=0.929 total time=
1.2s
[CV 4/5] END learning_rate=0.2, n_estimators=100;; score=0.972 total time=
1.0s
[CV 5/5] END learning_rate=0.2, n_estimators=100;; score=0.969 total time=
0.9s
[CV 1/5] END learning_rate=0.7, n_estimators=700;; score=0.925 total time=
2.4s
[CV 2/5] END learning_rate=0.7, n_estimators=700;; score=0.927 total time=
2.5s
[CV 3/5] END learning_rate=0.7, n_estimators=700;; score=0.899 total time=
2.2s
[CV 4/5] END learning_rate=0.7, n_estimators=700;; score=0.959 total time=
2.3s
[CV 5/5] END learning_rate=0.7, n_estimators=700;; score=0.931 total time=
3.0s
```



```
[CV 1/5] END learning_rate=0.4, n_estimators=1000;, score=0.927 total time
= 2.7s
[CV 2/5] END learning_rate=0.4, n_estimators=1000;, score=0.961 total time
= 3.5s
[CV 3/5] END learning_rate=0.4, n_estimators=1000;, score=0.884 total time
= 2.9s
[CV 4/5] END learning_rate=0.4, n_estimators=1000;, score=0.975 total time
= 3.5s
[CV 5/5] END learning_rate=0.4, n_estimators=1000;, score=0.941 total time
= 2.9s
[CV 1/5] END learning_rate=1, n_estimators=1000;, score=0.887 total time=
2.8s
[CV 2/5] END learning_rate=1, n_estimators=1000;, score=0.919 total time=
2.3s
[CV 3/5] END learning_rate=1, n_estimators=1000;, score=0.815 total time=
2.9s
[CV 4/5] END learning_rate=1, n_estimators=1000;, score=0.935 total time=
2.7s
[CV 5/5] END learning_rate=1, n_estimators=1000;, score=0.920 total time=
2.9s
[CV 1/5] END learning_rate=0.4, n_estimators=100;, score=0.927 total time=
0.9s
[CV 2/5] END learning_rate=0.4, n_estimators=100;, score=0.961 total time=
1.1s
[CV 3/5] END learning_rate=0.4, n_estimators=100;, score=0.884 total time=
1.0s
[CV 4/5] END learning_rate=0.4, n_estimators=100;, score=0.975 total time=
0.9s
[CV 5/5] END learning_rate=0.4, n_estimators=100;, score=0.941 total time=
0.9s
[CV 1/5] END .learning_rate=1, n_estimators=700;, score=0.887 total time=
2.1s
[CV 2/5] END .learning_rate=1, n_estimators=700;, score=0.919 total time=
2.1s
[CV 3/5] END .learning_rate=1, n_estimators=700;, score=0.815 total time=
2.2s
[CV 4/5] END .learning_rate=1, n_estimators=700;, score=0.935 total time=
2.3s
[CV 5/5] END .learning_rate=1, n_estimators=700;, score=0.920 total time=
2.3s
Mean Absolute Percentage Error: 0.04412554970177765
Root Mean Squared Error: 134966.0300207573
Root Mean Squared Log Error: 0.0697660248534803
R2 Score: 0.9556516116917408
```

In [123]:

```
grid_cat.best_params_
```

Out[123]:

```
{'n_estimators': 1000, 'learning_rate': 0.2}
```

In [124]:

```
param_grid = {'n_estimators': [100,300,500,800,1000],
              'max_samples': [0.24,0.58,0.71,0.96],
              'max_features': np.linspace(0,1,5),
              'bootstrap': [True,False],
              'oob_score': [True,False]
              }
grid_bag = RandomizedSearchCV(BaggingRegressor(),param_grid,verbose=2,cv=5)
train_and_evaluate_model(grid_bag)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[CV] END bootstrap=False, max_features=0.0, max_samples=0.58, n_estimators=800, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, max_features=0.0, max_samples=0.58, n_estimators=800, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, max_features=0.0, max_samples=0.58, n_estimators=800, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, max_features=0.0, max_samples=0.58, n_estimators=800, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, max_features=0.0, max_samples=0.58, n_estimators=800, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, max_features=0.5, max_samples=0.24, n_estimators=500, oob_score=False; total time= 2.0s
[CV] END bootstrap=False, max_features=0.5, max_samples=0.24, n_estimators=500, oob_score=False; total time= 1.8s
[CV] END bootstrap=False, max_features=0.5, max_samples=0.24, n_estimators=500, oob_score=False; total time= 1.9s
[CV] END bootstrap=False, max_features=0.5, max_samples=0.24, n_estimators=500, oob_score=False; total time= 1.8s
[CV] END bootstrap=False, max_features=0.5, max_samples=0.24, n_estimators=500, oob_score=False; total time= 1.8s
[CV] END bootstrap=False, max_features=1.0, max_samples=0.58, n_estimators=1000, oob_score=False; total time= 4.7s
[CV] END bootstrap=False, max_features=1.0, max_samples=0.58, n_estimators=1000, oob_score=False; total time= 4.7s
[CV] END bootstrap=False, max_features=1.0, max_samples=0.58, n_estimators=1000, oob_score=False; total time= 4.8s
[CV] END bootstrap=False, max_features=1.0, max_samples=0.58, n_estimators=1000, oob_score=False; total time= 4.7s
[CV] END bootstrap=False, max_features=1.0, max_samples=0.58, n_estimators=1000, oob_score=False; total time= 4.9s
[CV] END bootstrap=False, max_features=0.0, max_samples=0.58, n_estimators=500, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, max_features=0.0, max_samples=0.58, n_estimators=500, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, max_features=0.0, max_samples=0.58, n_estimators=500, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, max_features=0.0, max_samples=0.58, n_estimators=500, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, max_features=0.0, max_samples=0.58, n_estimators=500, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, max_features=0.75, max_samples=0.24, n_estimators=300, oob_score=False; total time= 1.3s
[CV] END bootstrap=False, max_features=0.75, max_samples=0.24, n_estimators=300, oob_score=False; total time= 1.2s
[CV] END bootstrap=False, max_features=0.75, max_samples=0.24, n_estimators=300, oob_score=False; total time= 1.1s
[CV] END bootstrap=False, max_features=0.75, max_samples=0.24, n_estimators=300, oob_score=False; total time= 1.1s
[CV] END bootstrap=False, max_features=0.75, max_samples=0.24, n_estimators=300, oob_score=False; total time= 1.1s
[CV] END bootstrap=False, max_features=0.75, max_samples=0.71, n_estimators=500, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, max_features=0.75, max_samples=0.71, n_estimators=500, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, max_features=0.75, max_samples=0.71, n_estimators=500, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, max_features=0.75, max_samples=0.71, n_estimators=500, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, max_features=0.75, max_samples=0.71, n_estimators=500, oob_score=True; total time= 0.0s
```

```
[CV] END bootstrap=True, max_features=0.0, max_samples=0.24, n_estimators=
800, oob_score=True; total time= 0.0s
[CV] END bootstrap=True, max_features=0.0, max_samples=0.24, n_estimators=
800, oob_score=True; total time= 0.0s
[CV] END bootstrap=True, max_features=0.0, max_samples=0.24, n_estimators=
800, oob_score=True; total time= 0.0s
[CV] END bootstrap=True, max_features=0.0, max_samples=0.24, n_estimators=
800, oob_score=True; total time= 0.0s
[CV] END bootstrap=True, max_features=0.0, max_samples=0.24, n_estimators=
800, oob_score=True; total time= 0.0s
[CV] END bootstrap=True, max_features=0.75, max_samples=0.58, n_estimators
=800, oob_score=False; total time= 3.5s
[CV] END bootstrap=True, max_features=0.75, max_samples=0.58, n_estimators
=800, oob_score=False; total time= 3.7s
[CV] END bootstrap=True, max_features=0.75, max_samples=0.58, n_estimators
=800, oob_score=False; total time= 3.6s
[CV] END bootstrap=True, max_features=0.75, max_samples=0.58, n_estimators
=800, oob_score=False; total time= 3.4s
[CV] END bootstrap=True, max_features=0.75, max_samples=0.58, n_estimators
=800, oob_score=False; total time= 3.4s
[CV] END bootstrap=True, max_features=0.5, max_samples=0.96, n_estimators=
500, oob_score=False; total time= 2.2s
[CV] END bootstrap=True, max_features=0.5, max_samples=0.96, n_estimators=
500, oob_score=False; total time= 2.2s
[CV] END bootstrap=True, max_features=0.5, max_samples=0.96, n_estimators=
500, oob_score=False; total time= 2.2s
[CV] END bootstrap=True, max_features=0.5, max_samples=0.96, n_estimators=
500, oob_score=False; total time= 2.1s
[CV] END bootstrap=True, max_features=0.5, max_samples=0.96, n_estimators=
500, oob_score=False; total time= 2.2s
[CV] END bootstrap=True, max_features=0.0, max_samples=0.58, n_estimators=
800, oob_score=False; total time= 0.0s
[CV] END bootstrap=True, max_features=0.0, max_samples=0.58, n_estimators=
800, oob_score=False; total time= 0.0s
[CV] END bootstrap=True, max_features=0.0, max_samples=0.58, n_estimators=
800, oob_score=False; total time= 0.0s
[CV] END bootstrap=True, max_features=0.0, max_samples=0.58, n_estimators=
800, oob_score=False; total time= 0.0s
[CV] END bootstrap=True, max_features=0.0, max_samples=0.58, n_estimators=
800, oob_score=False; total time= 0.0s
Mean Absolute Percentage Error: 0.031812137397857515
Root Mean Squared Error: 94700.60976805227
Root Mean Squared Log Error: 0.04551408853916175
R2 Score: 0.978165942716579
```

In [125]:

```
grid_bag.best_params_
```

Out[125]:

```
{'oob_score': False,
 'n_estimators': 1000,
 'max_samples': 0.58,
 'max_features': 1.0,
 'bootstrap': False}
```

In [126]:

```
param_grid = {'boosting_type': ['gbdt', 'dart', 'goss', 'rf'],
              'learning_rate': np.linspace(0,1,6)[1:],
              'n_estimators': [100,300,500,800,1000],
              'importance_type': ['split', 'gain'],
              'min_split_gain': [0.68,0.79,0.87,1]}

grid_lgbm = RandomizedSearchCV(LGBMRegressor(),param_grid,verbose=3)
train_and_evaluate_model(grid_lgbm)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[CV 1/5] END boosting_type=goss, importance_type=gain, learning_rate=1.0,
min_split_gain=1, n_estimators=300;; score=0.786 total time= 0.0s
[CV 2/5] END boosting_type=goss, importance_type=gain, learning_rate=1.0,
min_split_gain=1, n_estimators=300;; score=0.846 total time= 0.0s
[CV 3/5] END boosting_type=goss, importance_type=gain, learning_rate=1.0,
min_split_gain=1, n_estimators=300;; score=0.665 total time= 0.0s
[CV 4/5] END boosting_type=goss, importance_type=gain, learning_rate=1.0,
min_split_gain=1, n_estimators=300;; score=0.811 total time= 0.0s
[CV 5/5] END boosting_type=goss, importance_type=gain, learning_rate=1.0,
min_split_gain=1, n_estimators=300;; score=0.802 total time= 0.0s
[CV 1/5] END boosting_type=goss, importance_type=split, learning_rate=0.4,
min_split_gain=0.68, n_estimators=1000;; score=0.663 total time= 0.0s
[CV 2/5] END boosting_type=goss, importance_type=split, learning_rate=0.4,
min_split_gain=0.68, n_estimators=1000;; score=0.740 total time= 0.0s
[CV 3/5] END boosting_type=goss, importance_type=split, learning_rate=0.4,
min_split_gain=0.68, n_estimators=1000;; score=0.708 total time= 0.0s
[CV 4/5] END boosting_type=goss, importance_type=split, learning_rate=0.4,
min_split_gain=0.68, n_estimators=1000;; score=0.801 total time= 0.0s
[CV 5/5] END boosting_type=goss, importance_type=split, learning_rate=0.4,
min_split_gain=0.68, n_estimators=1000;; score=0.763 total time= 0.0s
[CV 1/5] END boosting_type=dart, importance_type=split, learning_rate=0.2,
min_split_gain=0.87, n_estimators=500;; score=0.931 total time= 0.2s
[CV 2/5] END boosting_type=dart, importance_type=split, learning_rate=0.2,
min_split_gain=0.87, n_estimators=500;; score=0.955 total time= 0.1s
[CV 3/5] END boosting_type=dart, importance_type=split, learning_rate=0.2,
min_split_gain=0.87, n_estimators=500;; score=0.953 total time= 0.1s
[CV 4/5] END boosting_type=dart, importance_type=split, learning_rate=0.2,
min_split_gain=0.87, n_estimators=500;; score=0.963 total time= 0.1s
[CV 5/5] END boosting_type=dart, importance_type=split, learning_rate=0.2,
min_split_gain=0.87, n_estimators=500;; score=0.964 total time= 0.2s
[CV 1/5] END boosting_type=gbdt, importance_type=gain, learning_rate=0.2,
min_split_gain=0.79, n_estimators=300;; score=0.940 total time= 0.0s
[CV 2/5] END boosting_type=gbdt, importance_type=gain, learning_rate=0.2,
min_split_gain=0.79, n_estimators=300;; score=0.957 total time= 0.0s
[CV 3/5] END boosting_type=gbdt, importance_type=gain, learning_rate=0.2,
min_split_gain=0.79, n_estimators=300;; score=0.942 total time= 0.0s
[CV 4/5] END boosting_type=gbdt, importance_type=gain, learning_rate=0.2,
min_split_gain=0.79, n_estimators=300;; score=0.955 total time= 0.0s
[CV 5/5] END boosting_type=gbdt, importance_type=gain, learning_rate=0.2,
min_split_gain=0.79, n_estimators=300;; score=0.957 total time= 0.0s
[CV 1/5] END boosting_type=gbdt, importance_type=split, learning_rate=1.0,
min_split_gain=0.79, n_estimators=500;; score=0.954 total time= 0.0s
[CV 2/5] END boosting_type=gbdt, importance_type=split, learning_rate=1.0,
min_split_gain=0.79, n_estimators=500;; score=0.953 total time= 0.0s
[CV 3/5] END boosting_type=gbdt, importance_type=split, learning_rate=1.0,
min_split_gain=0.79, n_estimators=500;; score=0.873 total time= 0.0s
[CV 4/5] END boosting_type=gbdt, importance_type=split, learning_rate=1.0,
min_split_gain=0.79, n_estimators=500;; score=0.898 total time= 0.0s
[CV 5/5] END boosting_type=gbdt, importance_type=split, learning_rate=1.0,
min_split_gain=0.79, n_estimators=500;; score=0.942 total time= 0.0s
[CV 1/5] END boosting_type=dart, importance_type=gain, learning_rate=1.0,
min_split_gain=0.79, n_estimators=800;; score=0.942 total time= 0.4s
[CV 2/5] END boosting_type=dart, importance_type=gain, learning_rate=1.0,
min_split_gain=0.79, n_estimators=800;; score=0.964 total time= 0.4s
[CV 3/5] END boosting_type=dart, importance_type=gain, learning_rate=1.0,
min_split_gain=0.79, n_estimators=800;; score=0.909 total time= 0.4s
[CV 4/5] END boosting_type=dart, importance_type=gain, learning_rate=1.0,
min_split_gain=0.79, n_estimators=800;; score=0.914 total time= 0.4s
[CV 5/5] END boosting_type=dart, importance_type=gain, learning_rate=1.0,
min_split_gain=0.79, n_estimators=800;; score=0.958 total time= 0.4s
```

```
[CV 1/5] END boosting_type=gbdt, importance_type=gain, learning_rate=0.4,
min_split_gain=0.87, n_estimators=100;; score=0.943 total time= 0.0s
[CV 2/5] END boosting_type=gbdt, importance_type=gain, learning_rate=0.4,
min_split_gain=0.87, n_estimators=100;; score=0.958 total time= 0.0s
[CV 3/5] END boosting_type=gbdt, importance_type=gain, learning_rate=0.4,
min_split_gain=0.87, n_estimators=100;; score=0.933 total time= 0.0s
[CV 4/5] END boosting_type=gbdt, importance_type=gain, learning_rate=0.4,
min_split_gain=0.87, n_estimators=100;; score=0.946 total time= 0.0s
[CV 5/5] END boosting_type=gbdt, importance_type=gain, learning_rate=0.4,
min_split_gain=0.87, n_estimators=100;; score=0.957 total time= 0.0s
[CV 1/5] END boosting_type=goss, importance_type=gain, learning_rate=0.4,
min_split_gain=0.79, n_estimators=100;; score=0.663 total time= 0.0s
[CV 2/5] END boosting_type=goss, importance_type=gain, learning_rate=0.4,
min_split_gain=0.79, n_estimators=100;; score=0.740 total time= 0.0s
[CV 3/5] END boosting_type=goss, importance_type=gain, learning_rate=0.4,
min_split_gain=0.79, n_estimators=100;; score=0.708 total time= 0.0s
[CV 4/5] END boosting_type=goss, importance_type=gain, learning_rate=0.4,
min_split_gain=0.79, n_estimators=100;; score=0.801 total time= 0.0s
[CV 5/5] END boosting_type=goss, importance_type=gain, learning_rate=0.4,
min_split_gain=0.79, n_estimators=100;; score=0.763 total time= 0.0s
[CV 1/5] END boosting_type=dart, importance_type=gain, learning_rate=0.8,
min_split_gain=1, n_estimators=300;; score=0.949 total time= 0.0s
[CV 2/5] END boosting_type=dart, importance_type=gain, learning_rate=0.8,
min_split_gain=1, n_estimators=300;; score=0.968 total time= 0.0s
[CV 3/5] END boosting_type=dart, importance_type=gain, learning_rate=0.8,
min_split_gain=1, n_estimators=300;; score=0.913 total time= 0.0s
[CV 4/5] END boosting_type=dart, importance_type=gain, learning_rate=0.8,
min_split_gain=1, n_estimators=300;; score=0.932 total time= 0.0s
[CV 5/5] END boosting_type=dart, importance_type=gain, learning_rate=0.8,
min_split_gain=1, n_estimators=300;; score=0.959 total time= 0.0s
[CV 1/5] END boosting_type=dart, importance_type=gain, learning_rate=1.0,
min_split_gain=0.87, n_estimators=500;; score=0.944 total time= 0.1s
[CV 2/5] END boosting_type=dart, importance_type=gain, learning_rate=1.0,
min_split_gain=0.87, n_estimators=500;; score=0.963 total time= 0.2s
[CV 3/5] END boosting_type=dart, importance_type=gain, learning_rate=1.0,
min_split_gain=0.87, n_estimators=500;; score=0.908 total time= 0.1s
[CV 4/5] END boosting_type=dart, importance_type=gain, learning_rate=1.0,
min_split_gain=0.87, n_estimators=500;; score=0.911 total time= 0.1s
[CV 5/5] END boosting_type=dart, importance_type=gain, learning_rate=1.0,
min_split_gain=0.87, n_estimators=500;; score=0.956 total time= 0.1s
Mean Absolute Percentage Error: 0.05287068945686181
Root Mean Squared Error: 142237.58139948634
Root Mean Squared Log Error: 0.07440089741866754
R2 Score: 0.9507441733056884
```

In [127]:

```
grid_lgbm.best_params_
```

Out[127]:

```
{'n_estimators': 500,
 'min_split_gain': 0.87,
 'learning_rate': 0.2,
 'importance_type': 'split',
 'boosting_type': 'dart'}
```

In [128]:

```
param_grid = {'n_estimators': [100,300,500,800,1000],
              'criterion': ["squared_error", "absolute_error", "friedman_mse", "poisson"],
              'max_features': ['auto','sqrt','log2'],
              'bootstrap': [True,False],
              'oob_score': [True,False],
              'max_samples': [0.2,0.45,0.7,0.95]
              }

grid_et = RandomizedSearchCV(ExtraTreesRegressor(),param_grid,verbose=2)
train_and_evaluate_model(grid_et)
```


Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[CV] END bootstrap=False, criterion=poisson, max_features=log2, max_sample
s=0.7, n_estimators=300, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, criterion=poisson, max_features=log2, max_sample
s=0.7, n_estimators=300, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, criterion=poisson, max_features=log2, max_sample
s=0.7, n_estimators=300, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, criterion=poisson, max_features=log2, max_sample
s=0.7, n_estimators=300, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, criterion=poisson, max_features=log2, max_sample
s=0.7, n_estimators=300, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, criterion=friedman_mse, max_features=log2, max_s
amples=0.45, n_estimators=300, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, criterion=friedman_mse, max_features=log2, max_s
amples=0.45, n_estimators=300, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, criterion=friedman_mse, max_features=log2, max_s
amples=0.45, n_estimators=300, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, criterion=friedman_mse, max_features=log2, max_s
amples=0.45, n_estimators=300, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, criterion=friedman_mse, max_features=log2, max_s
amples=0.45, n_estimators=300, oob_score=False; total time= 0.0s
[CV] END bootstrap=True, criterion=absolute_error, max_features=sqrt, max_
samples=0.7, n_estimators=1000, oob_score=True; total time= 4.6s
[CV] END bootstrap=True, criterion=absolute_error, max_features=sqrt, max_
samples=0.7, n_estimators=1000, oob_score=True; total time= 4.7s
[CV] END bootstrap=True, criterion=absolute_error, max_features=sqrt, max_
samples=0.7, n_estimators=1000, oob_score=True; total time= 4.7s
[CV] END bootstrap=True, criterion=absolute_error, max_features=sqrt, max_
samples=0.7, n_estimators=1000, oob_score=True; total time= 4.7s
[CV] END bootstrap=True, criterion=absolute_error, max_features=sqrt, max_
samples=0.7, n_estimators=1000, oob_score=True; total time= 4.7s
[CV] END bootstrap=True, criterion=squared_error, max_features=auto, max_s
amples=0.95, n_estimators=300, oob_score=False; total time= 1.0s
[CV] END bootstrap=True, criterion=squared_error, max_features=auto, max_s
amples=0.95, n_estimators=300, oob_score=False; total time= 1.1s
[CV] END bootstrap=True, criterion=squared_error, max_features=auto, max_s
amples=0.95, n_estimators=300, oob_score=False; total time= 1.2s
[CV] END bootstrap=True, criterion=squared_error, max_features=auto, max_s
amples=0.95, n_estimators=300, oob_score=False; total time= 1.1s
[CV] END bootstrap=True, criterion=squared_error, max_features=auto, max_s
amples=0.95, n_estimators=300, oob_score=False; total time= 1.1s
[CV] END bootstrap=False, criterion=absolute_error, max_features=sqrt, max
_samples=0.2, n_estimators=800, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, criterion=absolute_error, max_features=sqrt, max
_samples=0.2, n_estimators=800, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, criterion=absolute_error, max_features=sqrt, max
_samples=0.2, n_estimators=800, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, criterion=absolute_error, max_features=sqrt, max
_samples=0.2, n_estimators=800, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, criterion=absolute_error, max_features=sqrt, max
_samples=0.2, n_estimators=800, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, criterion=friedman_mse, max_features=sqrt, max_s
amples=0.7, n_estimators=300, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, criterion=friedman_mse, max_features=sqrt, max_s
amples=0.7, n_estimators=300, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, criterion=friedman_mse, max_features=sqrt, max_s
amples=0.7, n_estimators=300, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, criterion=friedman_mse, max_features=sqrt, max_s
amples=0.7, n_estimators=300, oob_score=True; total time= 0.0s
```

```
[CV] END bootstrap=False, criterion=absolute_error, max_features=sqrt, max_
_samples=0.95, n_estimators=500, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, criterion=absolute_error, max_features=sqrt, max_
_samples=0.95, n_estimators=500, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, criterion=absolute_error, max_features=sqrt, max_
_samples=0.95, n_estimators=500, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, criterion=absolute_error, max_features=sqrt, max_
_samples=0.95, n_estimators=500, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, criterion=absolute_error, max_features=sqrt, max_
_samples=0.95, n_estimators=500, oob_score=True; total time= 0.0s
[CV] END bootstrap=False, criterion=squared_error, max_features=log2, max_
_samples=0.7, n_estimators=500, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, criterion=squared_error, max_features=log2, max_
_samples=0.7, n_estimators=500, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, criterion=squared_error, max_features=log2, max_
_samples=0.7, n_estimators=500, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, criterion=squared_error, max_features=log2, max_
_samples=0.7, n_estimators=500, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, criterion=squared_error, max_features=log2, max_
_samples=0.7, n_estimators=500, oob_score=False; total time= 0.0s
[CV] END bootstrap=True, criterion=absolute_error, max_features=sqrt, max_
_samples=0.2, n_estimators=800, oob_score=True; total time= 3.4s
[CV] END bootstrap=True, criterion=absolute_error, max_features=sqrt, max_
_samples=0.2, n_estimators=800, oob_score=True; total time= 3.2s
[CV] END bootstrap=True, criterion=absolute_error, max_features=sqrt, max_
_samples=0.2, n_estimators=800, oob_score=True; total time= 3.1s
[CV] END bootstrap=True, criterion=absolute_error, max_features=sqrt, max_
_samples=0.2, n_estimators=800, oob_score=True; total time= 3.1s
[CV] END bootstrap=True, criterion=absolute_error, max_features=sqrt, max_
_samples=0.2, n_estimators=800, oob_score=True; total time= 3.1s
[CV] END bootstrap=False, criterion=absolute_error, max_features=auto, max_
_samples=0.45, n_estimators=100, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, criterion=absolute_error, max_features=auto, max_
_samples=0.45, n_estimators=100, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, criterion=absolute_error, max_features=auto, max_
_samples=0.45, n_estimators=100, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, criterion=absolute_error, max_features=auto, max_
_samples=0.45, n_estimators=100, oob_score=False; total time= 0.0s
[CV] END bootstrap=False, criterion=absolute_error, max_features=auto, max_
_samples=0.45, n_estimators=100, oob_score=False; total time= 0.0s
Mean Absolute Percentage Error: 0.024177768172682246
Root Mean Squared Error: 77496.62330560893
Root Mean Squared Log Error: 0.039356923308345025
R2 Score: 0.9853784149714498
```

In [129]:

```
grid_et.best_params_
```

Out[129]:

```
{'oob_score': False,
 'n_estimators': 300,
 'max_samples': 0.95,
 'max_features': 'auto',
 'criterion': 'squared_error',
 'bootstrap': True}
```

In [130]:

```

param_grid = {'n_estimators': [100,400,700,900,1000],
              'grow_policy': [0,1],
              'learning_rate': [0.1,0.4,0.6,0.8,1],
              'booster': ['gbtree','gblinear','dart'],
              'sampling_method': ['uniform','gradient_based'],
              'importance_type': ['gain','weight','cover','total_gain','total_cover']}

grid_xgb = RandomizedSearchCV(XGBRegressor(),param_grid,verbose=2,cv=5)
train_and_evaluate_model(grid_xgb)

```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```

[CV] END booster=gbtree, grow_policy=0, importance_type=gain, learning_
rate=0.6, n_estimators=900, sampling_method=uniform; total time=  0.0s
[CV] END booster=gbtree, grow_policy=0, importance_type=gain, learning_
rate=0.6, n_estimators=900, sampling_method=uniform; total time=  0.0s
[CV] END booster=gbtree, grow_policy=0, importance_type=gain, learning_
rate=0.6, n_estimators=900, sampling_method=uniform; total time=  0.0s
[CV] END booster=gbtree, grow_policy=0, importance_type=gain, learning_
rate=0.6, n_estimators=900, sampling_method=uniform; total time=  0.0s
[CV] END booster=gbtree, grow_policy=0, importance_type=gain, learning_
rate=0.6, n_estimators=900, sampling_method=uniform; total time=  0.0s
[CV] END booster=gbtree, grow_policy=0, importance_type=total_cover, le
arning_rate=0.4, n_estimators=700, sampling_method=uniform; total time=
0.0s
[CV] END booster=gbtree, grow_policy=0, importance_type=total_cover, le
arning_rate=0.4, n_estimators=700, sampling_method=uniform; total time=
0.0s
[CV] END booster=gbtree, grow_policy=0, importance_type=total_cover, le
arning_rate=0.4, n_estimators=700, sampling_method=uniform; total time=
^ ^

```

In [131]:

```
grid_xgb.best_params_
```

Out[131]:

```

{'sampling_method': 'uniform',
 'n_estimators': 1000,
 'learning_rate': 0.4,
 'importance_type': 'weight',
 'grow_policy': 0,
 'booster': 'gblinear'}

```

In [132]:

```

param_grid = {'loss': ['squared_error', 'huber', 'epsilon_insensitive', 'squared_epsilon_in
                'penalty': ['l2', 'l1', 'elasticnet'],
                'l1_ratio': [0.15, 0.45, 0.68, 0.81, 0.97],
                'alpha': [0.0001, 0.001, 0.01, 0.1, 1],
                'shuffle': [True, False],
                'learning_rate': ['adaptive', 'constant', 'optimal', 'invscaling'],
                'epsilon': np.linspace(0.001, 100, 10),
                'average': [True, False]
            }

grid_sgd = RandomizedSearchCV(SGDRegressor(), param_grid, verbose=2, cv=5)
train_and_evaluate_model(grid_sgd)

```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```

[CV] END alpha=0.0001, average=False, epsilon=77.77799999999999, l1_rat
io=0.81, learning_rate=invscaling, loss=squared_epsilon_insensitive, pe
nalty=elasticnet, shuffle=False; total time= 0.0s
[CV] END alpha=0.0001, average=False, epsilon=77.77799999999999, l1_rat
io=0.81, learning_rate=invscaling, loss=squared_epsilon_insensitive, pe
nalty=elasticnet, shuffle=False; total time= 0.0s
[CV] END alpha=0.0001, average=False, epsilon=77.77799999999999, l1_rat
io=0.81, learning_rate=invscaling, loss=squared_epsilon_insensitive, pe
nalty=elasticnet, shuffle=False; total time= 0.0s
[CV] END alpha=0.0001, average=False, epsilon=77.77799999999999, l1_rat
io=0.81, learning_rate=invscaling, loss=squared_epsilon_insensitive, pe
nalty=elasticnet, shuffle=False; total time= 0.0s
[CV] END alpha=0.0001, average=False, epsilon=77.77799999999999, l1_rat
io=0.81, learning_rate=invscaling, loss=squared_epsilon_insensitive, pe
nalty=elasticnet, shuffle=False; total time= 0.0s
[CV] END alpha=0.1, average=False, epsilon=33.333999999999996, l1_ratio
=0.15, learning_rate=optimal, loss=huber, penalty=elasticnet, shuffle=T
rue; total time= 0.0s

```

In [133]:

```
grid_sgd.best_params_
```

Out[133]:

```

{'shuffle': False,
 'penalty': 'elasticnet',
 'loss': 'squared_epsilon_insensitive',
 'learning_rate': 'invscaling',
 'l1_ratio': 0.81,
 'epsilon': 77.77799999999999,
 'average': False,
 'alpha': 0.0001}

```

In [134]:

```
param_grid = {'epsilon': np.linspace(1,10,10),  
              'alpha': np.linspace(0.0001,10,10)}  
  
grid_huber = RandomizedSearchCV(HuberRegressor(),param_grid,verbose=2,cv=5)  
train_and_evaluate_model(grid_huber)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[CV] END .....alpha=10.0, epsilon=10.0; total time=
0.0s
[CV] END .....alpha=10.0, epsilon=10.0; total time=
0.0s
[CV] END .....alpha=10.0, epsilon=10.0; total time=
0.0s
[CV] END .....alpha=10.0, epsilon=10.0; total time=
0.0s
[CV] END .....alpha=10.0, epsilon=10.0; total time=
0.0s
[CV] END .....alpha=7.777999999999999, epsilon=6.0; total time=
0.0s
[CV] END .....alpha=7.777999999999999, epsilon=6.0; total time=
0.0s
[CV] END .....alpha=7.777999999999999, epsilon=6.0; total time=
0.0s
[CV] END .....alpha=7.777999999999999, epsilon=6.0; total time=
0.0s
[CV] END .....alpha=7.777999999999999, epsilon=6.0; total time=
0.0s
[CV] END .....alpha=3.3334, epsilon=2.0; total time=
0.0s
[CV] END .....alpha=3.3334, epsilon=2.0; total time=
0.0s
[CV] END .....alpha=3.3334, epsilon=2.0; total time=
0.0s
[CV] END .....alpha=3.3334, epsilon=2.0; total time=
0.0s
[CV] END .....alpha=3.3334, epsilon=2.0; total time=
0.0s
[CV] END .....alpha=8.8889, epsilon=6.0; total time=
0.0s
[CV] END .....alpha=8.8889, epsilon=6.0; total time=
0.0s
[CV] END .....alpha=8.8889, epsilon=6.0; total time=
0.0s
[CV] END .....alpha=8.8889, epsilon=6.0; total time=
0.0s
[CV] END .....alpha=8.8889, epsilon=6.0; total time=
0.0s
[CV] END .....alpha=10.0, epsilon=1.0; total time=
0.0s
[CV] END .....alpha=10.0, epsilon=1.0; total time=
0.0s
[CV] END .....alpha=10.0, epsilon=1.0; total time=
0.0s
[CV] END .....alpha=10.0, epsilon=1.0; total time=
0.0s
[CV] END .....alpha=10.0, epsilon=1.0; total time=
0.0s
[CV] END .....alpha=4.4445, epsilon=8.0; total time=
0.0s
[CV] END .....alpha=4.4445, epsilon=8.0; total time=
0.0s
[CV] END .....alpha=4.4445, epsilon=8.0; total time=
0.0s
[CV] END .....alpha=4.4445, epsilon=8.0; total time=
0.0s
[CV] END .....alpha=4.4445, epsilon=8.0; total time=
0.0s
```

```
[CV] END .....alpha=8.8889, epsilon=5.0; total time=
0.0s
[CV] END .....alpha=8.8889, epsilon=5.0; total time=
0.0s
[CV] END .....alpha=8.8889, epsilon=5.0; total time=
0.0s
[CV] END .....alpha=8.8889, epsilon=5.0; total time=
0.0s
[CV] END .....alpha=8.8889, epsilon=5.0; total time=
0.0s
[CV] END .....alpha=10.0, epsilon=4.0; total time=
0.0s
[CV] END .....alpha=10.0, epsilon=4.0; total time=
0.0s
[CV] END .....alpha=10.0, epsilon=4.0; total time=
0.0s
[CV] END .....alpha=10.0, epsilon=4.0; total time=
0.0s
[CV] END .....alpha=10.0, epsilon=4.0; total time=
0.0s
[CV] END .....alpha=4.4445, epsilon=3.0; total time=
0.0s
[CV] END .....alpha=4.4445, epsilon=3.0; total time=
0.0s
[CV] END .....alpha=4.4445, epsilon=3.0; total time=
0.0s
[CV] END .....alpha=4.4445, epsilon=3.0; total time=
0.0s
[CV] END .....alpha=4.4445, epsilon=3.0; total time=
0.0s
[CV] END .....alpha=8.8889, epsilon=7.0; total time=
0.0s
[CV] END .....alpha=8.8889, epsilon=7.0; total time=
0.0s
[CV] END .....alpha=8.8889, epsilon=7.0; total time=
0.0s
[CV] END .....alpha=8.8889, epsilon=7.0; total time=
0.0s
[CV] END .....alpha=8.8889, epsilon=7.0; total time=
0.0s
Mean Absolute Percentage Error: 0.3882166185318009
Root Mean Squared Error: 640842.613070887
Root Mean Squared Log Error: 0.39938931342570616
R2 Score: 0.00015734670766076597
```

In [135]:

```
grid_huber.best_params_
```

Out[135]:

```
{'epsilon': 2.0, 'alpha': 3.3334}
```

In [136]:

```
param_grid = {'link': ['auto', 'identity', 'log'],
              'solver': ['lbfgs', 'newton-cholesky'],
              'power': [0,1,(1,2),2,3],
              'alpha': np.linspace(0.0001,10,10)}

grid_tweedie = RandomizedSearchCV(TweedieRegressor(),param_grid,verbose=2,cv=5)
train_and_evaluate_model(grid_tweedie)
```



```
Fitting 5 folds for each of 10 candidates, totalling 50 fits
[CV] END alpha=10.0, link=identity, power=(1, 2), solver=newton-cholesky;
total time= 0.0s
[CV] END alpha=10.0, link=identity, power=(1, 2), solver=newton-cholesky;
total time= 0.0s
[CV] END alpha=10.0, link=identity, power=(1, 2), solver=newton-cholesky;
total time= 0.0s
[CV] END alpha=10.0, link=identity, power=(1, 2), solver=newton-cholesky;
total time= 0.0s
[CV] END alpha=10.0, link=identity, power=(1, 2), solver=newton-cholesky;
total time= 0.0s
[CV] END alpha=2.2223, link=auto, power=3, solver=newton-cholesky; total t
ime= 0.0s
[CV] END alpha=2.2223, link=auto, power=3, solver=newton-cholesky; total t
ime= 0.0s
[CV] END alpha=2.2223, link=auto, power=3, solver=newton-cholesky; total t
ime= 0.0s
[CV] END alpha=2.2223, link=auto, power=3, solver=newton-cholesky; total t
ime= 0.0s
[CV] END alpha=2.2223, link=auto, power=3, solver=newton-cholesky; total t
ime= 0.0s
[CV] END alpha=2.2223, link=auto, power=3, solver=newton-cholesky; total t
ime= 0.0s
[CV] END alpha=5.5556, link=log, power=2, solver=newton-cholesky; total ti
me= 0.0s
[CV] END alpha=5.5556, link=log, power=2, solver=newton-cholesky; total ti
me= 0.0s
[CV] END alpha=5.5556, link=log, power=2, solver=newton-cholesky; total ti
me= 0.0s
[CV] END alpha=5.5556, link=log, power=2, solver=newton-cholesky; total ti
me= 0.0s
[CV] END alpha=5.5556, link=log, power=2, solver=newton-cholesky; total ti
me= 0.0s
[CV] END alpha=7.777999999999999, link=identity, power=(1, 2), solver=lbfg
s; total time= 0.0s
[CV] END alpha=7.777999999999999, link=identity, power=(1, 2), solver=lbfg
s; total time= 0.0s
[CV] END alpha=7.777999999999999, link=identity, power=(1, 2), solver=lbfg
s; total time= 0.0s
[CV] END alpha=7.777999999999999, link=identity, power=(1, 2), solver=lbfg
s; total time= 0.0s
[CV] END alpha=7.777999999999999, link=identity, power=(1, 2), solver=lbfg
s; total time= 0.0s
[CV] END alpha=7.777999999999999, link=identity, power=(1, 2), solver=lbfg
s; total time= 0.0s
[CV] END alpha=4.4445, link=log, power=(1, 2), solver=newton-cholesky; tot
al time= 0.0s
[CV] END alpha=4.4445, link=log, power=(1, 2), solver=newton-cholesky; tot
al time= 0.0s
[CV] END alpha=4.4445, link=log, power=(1, 2), solver=newton-cholesky; tot
al time= 0.0s
[CV] END alpha=4.4445, link=log, power=(1, 2), solver=newton-cholesky; tot
al time= 0.0s
[CV] END alpha=4.4445, link=log, power=(1, 2), solver=newton-cholesky; tot
al time= 0.0s
[CV] END alpha=4.4445, link=log, power=(1, 2), solver=newton-cholesky; tot
al time= 0.0s
[CV] END alpha=2.2223, link=auto, power=(1, 2), solver=lbfgs; total time=
0.0s
[CV] END alpha=2.2223, link=auto, power=(1, 2), solver=lbfgs; total time=
0.0s
[CV] END alpha=2.2223, link=auto, power=(1, 2), solver=lbfgs; total time=
0.0s
[CV] END alpha=2.2223, link=auto, power=(1, 2), solver=lbfgs; total time=
0.0s
[CV] END alpha=2.2223, link=auto, power=(1, 2), solver=lbfgs; total time=
0.0s
```

```
[CV] END alpha=6.6667, link=identity, power=3, solver=newton-cholesky; total time= 0.0s
[CV] END alpha=6.6667, link=identity, power=3, solver=newton-cholesky; total time= 0.0s
[CV] END alpha=6.6667, link=identity, power=3, solver=newton-cholesky; total time= 0.0s
[CV] END alpha=6.6667, link=identity, power=3, solver=newton-cholesky; total time= 0.0s
[CV] END alpha=6.6667, link=identity, power=3, solver=newton-cholesky; total time= 0.0s
[CV] END alpha=7.777999999999999, link=identity, power=3, solver=newton-cholesky; total time= 0.0s
[CV] END alpha=7.777999999999999, link=identity, power=3, solver=newton-cholesky; total time= 0.0s
[CV] END alpha=7.777999999999999, link=identity, power=3, solver=newton-cholesky; total time= 0.0s
[CV] END alpha=7.777999999999999, link=identity, power=3, solver=newton-cholesky; total time= 0.0s
[CV] END alpha=7.777999999999999, link=identity, power=3, solver=newton-cholesky; total time= 0.0s
[CV] END .....alpha=0.0001, link=auto, power=1, solver=lbfgs; total time= 0.0s
[CV] END .....alpha=0.0001, link=auto, power=1, solver=lbfgs; total time= 0.0s
[CV] END .....alpha=0.0001, link=auto, power=1, solver=lbfgs; total time= 0.0s
[CV] END .....alpha=0.0001, link=auto, power=1, solver=lbfgs; total time= 0.0s
[CV] END .....alpha=0.0001, link=auto, power=1, solver=lbfgs; total time= 0.0s
[CV] END alpha=0.0001, link=auto, power=2, solver=newton-cholesky; total time= 0.0s
[CV] END alpha=0.0001, link=auto, power=2, solver=newton-cholesky; total time= 0.0s
[CV] END alpha=0.0001, link=auto, power=2, solver=newton-cholesky; total time= 0.0s
[CV] END alpha=0.0001, link=auto, power=2, solver=newton-cholesky; total time= 0.0s
[CV] END alpha=0.0001, link=auto, power=2, solver=newton-cholesky; total time= 0.0s
Mean Absolute Percentage Error: 0.0591406631966941
Root Mean Squared Error: 136080.82844277137
Root Mean Squared Log Error: 0.07777764320505613
R2 Score: 0.9549159644382277
```

In [137]:

```
grid_tweedie.best_params_
```

Out[137]:

```
{'solver': 'newton-cholesky', 'power': 2, 'link': 'auto', 'alpha': 0.0001}
```

In [138]:

```
param_grid = {'C': [0.0001,0.001,0.01,0.1,1,10],
              'loss': ['epsilon_insensitive','squared_epsilon_insensitive'],
              'epsilon': np.linspace(0.001,1,5),
              'shuffle': [True,False],
              'average': [True,False]}

grid_pa = RandomizedSearchCV(PassiveAggressiveRegressor(),param_grid,verbose=2,cv=5)
train_and_evaluate_model(grid_pa)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[CV] END C=0.1, average=False, epsilon=0.25075, loss=squared_epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=0.1, average=False, epsilon=0.25075, loss=squared_epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=0.1, average=False, epsilon=0.25075, loss=squared_epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=0.1, average=False, epsilon=0.25075, loss=squared_epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=0.1, average=False, epsilon=0.25075, loss=squared_epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=0.0001, average=True, epsilon=0.75025, loss=epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=0.0001, average=True, epsilon=0.75025, loss=epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=0.0001, average=True, epsilon=0.75025, loss=epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=0.0001, average=True, epsilon=0.75025, loss=epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=0.0001, average=True, epsilon=0.75025, loss=epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=0.0001, average=True, epsilon=0.75025, loss=epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=1, average=True, epsilon=0.75025, loss=epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=1, average=True, epsilon=0.75025, loss=epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=1, average=True, epsilon=0.75025, loss=epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=1, average=True, epsilon=0.75025, loss=epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=1, average=True, epsilon=0.75025, loss=epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=10, average=True, epsilon=0.25075, loss=epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=10, average=True, epsilon=0.25075, loss=epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=10, average=True, epsilon=0.25075, loss=epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=10, average=True, epsilon=0.25075, loss=epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=10, average=True, epsilon=0.25075, loss=epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=10, average=True, epsilon=0.25075, loss=epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=0.0001, average=True, epsilon=0.25075, loss=squared_epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=0.0001, average=True, epsilon=0.25075, loss=squared_epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=0.0001, average=True, epsilon=0.25075, loss=squared_epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=0.0001, average=True, epsilon=0.25075, loss=squared_epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=0.0001, average=True, epsilon=0.25075, loss=squared_epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=10, average=True, epsilon=0.001, loss=squared_epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=10, average=True, epsilon=0.001, loss=squared_epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=10, average=True, epsilon=0.001, loss=squared_epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=10, average=True, epsilon=0.001, loss=squared_epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=10, average=True, epsilon=0.001, loss=squared_epsilon_insensitive, shuffle=True; total time= 0.0s
```

```
[CV] END C=1, average=True, epsilon=1.0, loss=epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=1, average=True, epsilon=1.0, loss=epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=1, average=True, epsilon=1.0, loss=epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=1, average=True, epsilon=1.0, loss=epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=1, average=True, epsilon=1.0, loss=epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=0.01, average=False, epsilon=0.75025, loss=epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=0.01, average=False, epsilon=0.75025, loss=epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=0.01, average=False, epsilon=0.75025, loss=epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=0.01, average=False, epsilon=0.75025, loss=epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=0.01, average=False, epsilon=0.75025, loss=epsilon_insensitive, shuffle=False; total time= 0.0s
[CV] END C=1, average=False, epsilon=0.25075, loss=squared_epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=1, average=False, epsilon=0.25075, loss=squared_epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=1, average=False, epsilon=0.25075, loss=squared_epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=1, average=False, epsilon=0.25075, loss=squared_epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=1, average=False, epsilon=0.25075, loss=squared_epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=0.001, average=True, epsilon=0.5005, loss=epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=0.001, average=True, epsilon=0.5005, loss=epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=0.001, average=True, epsilon=0.5005, loss=epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=0.001, average=True, epsilon=0.5005, loss=epsilon_insensitive, shuffle=True; total time= 0.0s
[CV] END C=0.001, average=True, epsilon=0.5005, loss=epsilon_insensitive, shuffle=True; total time= 0.0s
Mean Absolute Percentage Error: 0.0341132425926999
Root Mean Squared Error: 103082.83496662248
Root Mean Squared Log Error: 0.05185352883288553
R2 Score: 0.9741296918089055
```

In [139]:

```
grid_pa.best_params_
```

Out[139]:

```
{'shuffle': True,
 'loss': 'squared_epsilon_insensitive',
 'epsilon': 0.25075,
 'average': False,
 'C': 0.1}
```

In [140]:

```
param_grid = {  
    'eps': [0.0001, 0.001, 0.01, 0.1, 1],  
    'positive': [True, False],  
    'selection': ['cyclic', 'random']  
}  
  
grid_lasso = RandomizedSearchCV(LassoCV(), param_grid, verbose=2, cv=5)  
train_and_evaluate_model(grid_lasso)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[CV] END .....eps=0.001, positive=False, selection=cyclic; total time=
0.0s
[CV] END .....eps=0.001, positive=False, selection=cyclic; total time=
0.0s
[CV] END .....eps=0.001, positive=False, selection=cyclic; total time=
0.0s
[CV] END .....eps=0.001, positive=False, selection=cyclic; total time=
0.0s
[CV] END .....eps=0.001, positive=False, selection=cyclic; total time=
0.0s
[CV] END .....eps=1, positive=True, selection=random; total time=
0.0s
[CV] END .....eps=1, positive=True, selection=random; total time=
0.0s
[CV] END .....eps=1, positive=True, selection=random; total time=
0.0s
[CV] END .....eps=1, positive=True, selection=random; total time=
0.0s
[CV] END .....eps=1, positive=True, selection=random; total time=
0.0s
[CV] END .....eps=1, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....eps=1, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....eps=1, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....eps=1, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....eps=1, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....eps=0.1, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....eps=0.1, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....eps=0.1, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....eps=0.1, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....eps=0.1, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....eps=0.0001, positive=False, selection=cyclic; total time=
0.0s
[CV] END .....eps=0.0001, positive=False, selection=cyclic; total time=
0.0s
[CV] END .....eps=0.0001, positive=False, selection=cyclic; total time=
0.0s
[CV] END .....eps=0.0001, positive=False, selection=cyclic; total time=
0.0s
[CV] END .....eps=0.0001, positive=False, selection=cyclic; total time=
0.0s
[CV] END .....eps=0.0001, positive=False, selection=random; total time=
0.0s
[CV] END .....eps=0.0001, positive=False, selection=random; total time=
0.0s
[CV] END .....eps=0.0001, positive=False, selection=random; total time=
0.0s
[CV] END .....eps=0.0001, positive=False, selection=random; total time=
0.0s
[CV] END .....eps=0.0001, positive=False, selection=random; total time=
0.1s
```

```
[CV] END .....eps=0.001, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....eps=0.001, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....eps=0.001, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....eps=0.001, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....eps=0.001, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....eps=0.01, positive=True, selection=random; total time=
0.0s
[CV] END .....eps=0.01, positive=True, selection=random; total time=
0.0s
[CV] END .....eps=0.01, positive=True, selection=random; total time=
0.0s
[CV] END .....eps=0.01, positive=True, selection=random; total time=
0.0s
[CV] END .....eps=0.01, positive=True, selection=random; total time=
0.0s
[CV] END .....eps=0.001, positive=True, selection=random; total time=
0.0s
[CV] END .....eps=0.001, positive=True, selection=random; total time=
0.0s
[CV] END .....eps=0.001, positive=True, selection=random; total time=
0.0s
[CV] END .....eps=0.001, positive=True, selection=random; total time=
0.0s
[CV] END .....eps=0.001, positive=True, selection=random; total time=
0.0s
[CV] END .....eps=0.1, positive=True, selection=random; total time=
0.0s
[CV] END .....eps=0.1, positive=True, selection=random; total time=
0.0s
[CV] END .....eps=0.1, positive=True, selection=random; total time=
0.0s
[CV] END .....eps=0.1, positive=True, selection=random; total time=
0.0s
[CV] END .....eps=0.1, positive=True, selection=random; total time=
0.0s
Mean Absolute Percentage Error: 0.03607362447546001
Root Mean Squared Error: 89480.54698845446
Root Mean Squared Log Error: 0.052034517531026246
R2 Score: 0.9805066646085272
```

In [141]:

```
grid_lasso.best_params_
```

Out[141]:

```
{'selection': 'random', 'positive': True, 'eps': 0.01}
```


In [142]:

```
param_grid = {  
    'alphas': [(0.1, 1.0, 10.0), (0.01, 0.1, 1), (0.001, 0.01, 0.1)],  
    'gcv_mode': ['auto', 'svd', 'eigen'],  
    'store_cv_values': [True, False],  
    'alpha_per_target': [True, False]  
}  
  
grid_ridge = RandomizedSearchCV(RidgeCV(), param_grid, verbose=2, cv=5)  
train_and_evaluate_model(grid_ridge)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[CV] END alpha_per_target=True, alphas=(0.001, 0.01, 0.1), gcv_mode=svd, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.001, 0.01, 0.1), gcv_mode=svd, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.001, 0.01, 0.1), gcv_mode=svd, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.001, 0.01, 0.1), gcv_mode=svd, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.001, 0.01, 0.1), gcv_mode=svd, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=False, alphas=(0.001, 0.01, 0.1), gcv_mode=svd, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=False, alphas=(0.001, 0.01, 0.1), gcv_mode=svd, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=False, alphas=(0.001, 0.01, 0.1), gcv_mode=svd, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=False, alphas=(0.001, 0.01, 0.1), gcv_mode=svd, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=False, alphas=(0.001, 0.01, 0.1), gcv_mode=svd, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=False, alphas=(0.1, 1.0, 10.0), gcv_mode=eigen, store_cv_values=True; total time= 0.0s
[CV] END alpha_per_target=False, alphas=(0.1, 1.0, 10.0), gcv_mode=eigen, store_cv_values=True; total time= 0.0s
[CV] END alpha_per_target=False, alphas=(0.1, 1.0, 10.0), gcv_mode=eigen, store_cv_values=True; total time= 0.0s
[CV] END alpha_per_target=False, alphas=(0.1, 1.0, 10.0), gcv_mode=eigen, store_cv_values=True; total time= 0.0s
[CV] END alpha_per_target=False, alphas=(0.1, 1.0, 10.0), gcv_mode=eigen, store_cv_values=True; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.1, 1.0, 10.0), gcv_mode=auto, store_cv_values=True; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.1, 1.0, 10.0), gcv_mode=auto, store_cv_values=True; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.1, 1.0, 10.0), gcv_mode=auto, store_cv_values=True; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.1, 1.0, 10.0), gcv_mode=auto, store_cv_values=True; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.1, 1.0, 10.0), gcv_mode=auto, store_cv_values=True; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.01, 0.1, 1), gcv_mode=eigen, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.01, 0.1, 1), gcv_mode=eigen, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.01, 0.1, 1), gcv_mode=eigen, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.01, 0.1, 1), gcv_mode=eigen, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.01, 0.1, 1), gcv_mode=eigen, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=False, alphas=(0.01, 0.1, 1), gcv_mode=eigen, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=False, alphas=(0.01, 0.1, 1), gcv_mode=eigen, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=False, alphas=(0.01, 0.1, 1), gcv_mode=eigen, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=False, alphas=(0.01, 0.1, 1), gcv_mode=eigen, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=False, alphas=(0.01, 0.1, 1), gcv_mode=eigen, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=False, alphas=(0.01, 0.1, 1), gcv_mode=eigen, store_cv_values=False; total time= 0.0s
```

```
[CV] END alpha_per_target=True, alphas=(0.1, 1.0, 10.0), gcv_mode=eigen, store_cv_values=True; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.1, 1.0, 10.0), gcv_mode=eigen, store_cv_values=True; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.1, 1.0, 10.0), gcv_mode=eigen, store_cv_values=True; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.1, 1.0, 10.0), gcv_mode=eigen, store_cv_values=True; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.1, 1.0, 10.0), gcv_mode=eigen, store_cv_values=True; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.1, 1.0, 10.0), gcv_mode=eigen, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.1, 1.0, 10.0), gcv_mode=eigen, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.1, 1.0, 10.0), gcv_mode=eigen, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.1, 1.0, 10.0), gcv_mode=eigen, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.1, 1.0, 10.0), gcv_mode=eigen, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.01, 0.1, 1), gcv_mode=auto, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.01, 0.1, 1), gcv_mode=auto, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.01, 0.1, 1), gcv_mode=auto, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.01, 0.1, 1), gcv_mode=auto, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=True, alphas=(0.01, 0.1, 1), gcv_mode=auto, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=False, alphas=(0.01, 0.1, 1), gcv_mode=svd, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=False, alphas=(0.01, 0.1, 1), gcv_mode=svd, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=False, alphas=(0.01, 0.1, 1), gcv_mode=svd, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=False, alphas=(0.01, 0.1, 1), gcv_mode=svd, store_cv_values=False; total time= 0.0s
[CV] END alpha_per_target=False, alphas=(0.01, 0.1, 1), gcv_mode=svd, store_cv_values=False; total time= 0.0s
Mean Absolute Percentage Error: 0.03757672714457013
Root Mean Squared Error: 95459.32132352471
Root Mean Squared Log Error: 0.053988205973587854
R2 Score: 0.977814686026575
```

In [143]:

```
grid_ridge.best_params_
```

Out[143]:

```
{'store_cv_values': True,
 'gcv_mode': 'eigen',
 'alphas': (0.1, 1.0, 10.0),
 'alpha_per_target': False}
```

In [144]:

```
param_grid = {'selection': ['cyclic', 'random'],
              'l1_ratio': [0.1, 0.3, 0.5, 0.8, 1],
              'positive': [True, False]}

grid_elasticnet = RandomizedSearchCV(ElasticNetCV(), param_grid, verbose=2, cv=5)
train_and_evaluate_model(grid_elasticnet)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[CV] END .....l1_ratio=0.3, positive=True, selection=random; total time=0.0s
[CV] END .....l1_ratio=0.3, positive=True, selection=random; total time=0.0s
[CV] END .....l1_ratio=0.3, positive=True, selection=random; total time=0.0s
[CV] END .....l1_ratio=0.3, positive=True, selection=random; total time=0.0s
[CV] END .....l1_ratio=0.3, positive=True, selection=random; total time=0.0s
[CV] END .....l1_ratio=0.3, positive=False, selection=random; total time=0.0s
[CV] END .....l1_ratio=0.3, positive=False, selection=random; total time=0.0s
[CV] END .....l1_ratio=0.3, positive=False, selection=random; total time=0.0s
[CV] END .....l1_ratio=0.3, positive=False, selection=random; total time=0.0s
[CV] END .....l1_ratio=0.3, positive=False, selection=random; total time=0.0s
[CV] END .....l1_ratio=0.3, positive=False, selection=random; total time=0.0s
[CV] END .....l1_ratio=0.1, positive=False, selection=random; total time=0.0s
[CV] END .....l1_ratio=0.1, positive=False, selection=random; total time=0.0s
[CV] END .....l1_ratio=0.1, positive=False, selection=random; total time=0.0s
[CV] END .....l1_ratio=0.1, positive=False, selection=random; total time=0.0s
[CV] END .....l1_ratio=0.1, positive=False, selection=random; total time=0.0s
[CV] END .....l1_ratio=0.3, positive=False, selection=cyclic; total time=0.0s
[CV] END .....l1_ratio=0.3, positive=False, selection=cyclic; total time=0.0s
[CV] END .....l1_ratio=0.3, positive=False, selection=cyclic; total time=0.0s
[CV] END .....l1_ratio=0.3, positive=False, selection=cyclic; total time=0.0s
[CV] END .....l1_ratio=0.3, positive=False, selection=cyclic; total time=0.0s
[CV] END .....l1_ratio=0.5, positive=False, selection=random; total time=0.0s
[CV] END .....l1_ratio=0.5, positive=False, selection=random; total time=0.0s
[CV] END .....l1_ratio=0.5, positive=False, selection=random; total time=0.0s
[CV] END .....l1_ratio=0.5, positive=False, selection=random; total time=0.0s
[CV] END .....l1_ratio=0.5, positive=False, selection=random; total time=0.0s
[CV] END .....l1_ratio=1, positive=False, selection=cyclic; total time=0.0s
[CV] END .....l1_ratio=1, positive=False, selection=cyclic; total time=0.0s
[CV] END .....l1_ratio=1, positive=False, selection=cyclic; total time=0.0s
[CV] END .....l1_ratio=1, positive=False, selection=cyclic; total time=0.0s
```

```
[CV] END .....l1_ratio=1, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....l1_ratio=1, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....l1_ratio=1, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....l1_ratio=1, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....l1_ratio=1, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....l1_ratio=1, positive=False, selection=random; total time=
0.0s
[CV] END .....l1_ratio=1, positive=False, selection=random; total time=
0.0s
[CV] END .....l1_ratio=1, positive=False, selection=random; total time=
0.0s
[CV] END .....l1_ratio=1, positive=False, selection=random; total time=
0.0s
[CV] END .....l1_ratio=1, positive=False, selection=random; total time=
0.0s
[CV] END .....l1_ratio=0.3, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....l1_ratio=0.3, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....l1_ratio=0.3, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....l1_ratio=0.3, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....l1_ratio=0.3, positive=True, selection=cyclic; total time=
0.0s
[CV] END .....l1_ratio=0.8, positive=False, selection=random; total time=
0.0s
[CV] END .....l1_ratio=0.8, positive=False, selection=random; total time=
0.0s
[CV] END .....l1_ratio=0.8, positive=False, selection=random; total time=
0.0s
[CV] END .....l1_ratio=0.8, positive=False, selection=random; total time=
0.0s
[CV] END .....l1_ratio=0.8, positive=False, selection=random; total time=
0.0s
Mean Absolute Percentage Error: 0.03754926308061059
Root Mean Squared Error: 89235.81748679242
Root Mean Squared Log Error: 0.05393519130217834
R2 Score: 0.9806131474279399
```

In [145]:

```
grid_elasticnet.best_params_
```

Out[145]:

```
{'selection': 'cyclic', 'positive': True, 'l1_ratio': 1}
```

In [146]:

```
param_grid = {'n_estimators': np.arange(100,1001,300),  
              'learning_rate': [0.1,0.4,0.6,0.8,1],  
              'loss': ['linear', 'square', 'exponential']}  
  
grid_ab = RandomizedSearchCV(AdaBoostRegressor(),param_grid,cv=5,verbose=2)  
train_and_evaluate_model(grid_ab)
```

```
Fitting 5 folds for each of 10 candidates, totalling 50 fits
[CV] END ..learning_rate=0.8, loss=linear, n_estimators=1000; total time=
3.4s
[CV] END ..learning_rate=0.8, loss=linear, n_estimators=1000; total time=
3.3s
[CV] END ..learning_rate=0.8, loss=linear, n_estimators=1000; total time=
3.4s
[CV] END ..learning_rate=0.8, loss=linear, n_estimators=1000; total time=
3.4s
[CV] END ..learning_rate=0.8, loss=linear, n_estimators=1000; total time=
3.4s
[CV] END ...learning_rate=0.8, loss=square, n_estimators=100; total time=
0.2s
[CV] END ...learning_rate=0.8, loss=square, n_estimators=100; total time=
0.2s
[CV] END ...learning_rate=0.8, loss=square, n_estimators=100; total time=
0.3s
[CV] END ...learning_rate=0.8, loss=square, n_estimators=100; total time=
0.3s
[CV] END ...learning_rate=0.8, loss=square, n_estimators=100; total time=
0.2s
[CV] END learning_rate=0.6, loss=exponential, n_estimators=1000; total tim
e= 3.5s
[CV] END learning_rate=0.6, loss=exponential, n_estimators=1000; total tim
e= 3.8s
[CV] END learning_rate=0.6, loss=exponential, n_estimators=1000; total tim
e= 3.6s
[CV] END learning_rate=0.6, loss=exponential, n_estimators=1000; total tim
e= 3.5s
[CV] END learning_rate=0.6, loss=exponential, n_estimators=1000; total tim
e= 3.4s
[CV] END learning_rate=0.1, loss=exponential, n_estimators=1000; total tim
e= 3.4s
[CV] END learning_rate=0.1, loss=exponential, n_estimators=1000; total tim
e= 3.4s
[CV] END learning_rate=0.1, loss=exponential, n_estimators=1000; total tim
e= 3.4s
[CV] END learning_rate=0.1, loss=exponential, n_estimators=1000; total tim
e= 3.4s
[CV] END learning_rate=0.1, loss=exponential, n_estimators=1000; total tim
e= 3.4s
[CV] END learning_rate=0.1, loss=exponential, n_estimators=1000; total tim
e= 3.4s
[CV] END ...learning_rate=0.8, loss=square, n_estimators=700; total time=
2.2s
[CV] END ...learning_rate=0.8, loss=square, n_estimators=700; total time=
2.3s
[CV] END ...learning_rate=0.8, loss=square, n_estimators=700; total time=
2.3s
[CV] END ...learning_rate=0.8, loss=square, n_estimators=700; total time=
2.3s
[CV] END ...learning_rate=0.8, loss=square, n_estimators=700; total time=
2.5s
[CV] END learning_rate=0.8, loss=exponential, n_estimators=100; total time
= 0.2s
[CV] END learning_rate=0.8, loss=exponential, n_estimators=100; total time
= 0.3s
[CV] END learning_rate=0.8, loss=exponential, n_estimators=100; total time
= 0.3s
[CV] END learning_rate=0.8, loss=exponential, n_estimators=100; total time
= 0.3s
[CV] END learning_rate=0.8, loss=exponential, n_estimators=100; total time
= 0.3s
```



```
[CV] END ...learning_rate=0.8, loss=linear, n_estimators=400; total time=1.3s
[CV] END ...learning_rate=0.8, loss=linear, n_estimators=400; total time=1.3s
[CV] END ...learning_rate=0.8, loss=linear, n_estimators=400; total time=1.3s
[CV] END ...learning_rate=0.8, loss=linear, n_estimators=400; total time=1.3s
[CV] END ...learning_rate=0.8, loss=linear, n_estimators=400; total time=1.3s
[CV] END ...learning_rate=0.8, loss=linear, n_estimators=700; total time=2.3s
[CV] END ...learning_rate=0.8, loss=linear, n_estimators=700; total time=2.3s
[CV] END ...learning_rate=0.8, loss=linear, n_estimators=700; total time=2.3s
[CV] END ...learning_rate=0.8, loss=linear, n_estimators=700; total time=2.3s
[CV] END ...learning_rate=0.8, loss=linear, n_estimators=700; total time=2.3s
[CV] END ...learning_rate=0.8, loss=linear, n_estimators=700; total time=2.3s
[CV] END .....learning_rate=1, loss=linear, n_estimators=700; total time=2.3s
[CV] END .....learning_rate=1, loss=linear, n_estimators=700; total time=2.4s
[CV] END .....learning_rate=1, loss=linear, n_estimators=700; total time=2.6s
[CV] END .....learning_rate=1, loss=linear, n_estimators=700; total time=2.4s
[CV] END .....learning_rate=1, loss=linear, n_estimators=700; total time=2.3s
[CV] END ...learning_rate=0.8, loss=square, n_estimators=400; total time=1.4s
[CV] END ...learning_rate=0.8, loss=square, n_estimators=400; total time=1.5s
[CV] END ...learning_rate=0.8, loss=square, n_estimators=400; total time=1.4s
[CV] END ...learning_rate=0.8, loss=square, n_estimators=400; total time=1.3s
[CV] END ...learning_rate=0.8, loss=square, n_estimators=400; total time=1.3s
Mean Absolute Percentage Error: 0.06382350157680054
Root Mean Squared Error: 125385.84798589497
Root Mean Squared Log Error: 0.0759745650968994
R2 Score: 0.9617240536246201
```

In [147]:

```
grid_ab.best_params_
```

Out[147]:

```
{'n_estimators': 400, 'loss': 'square', 'learning_rate': 0.8}
```

In [148]:

```
param_grid = {'solver': ['lbfgs', 'newton-cholesky'],
              'fit_intercept': [True, False],
              'alpha': [0.001, 0.01, 0.1, 1, 10],
              'warm_start': [True, False]}

grid_poisson = RandomizedSearchCV(PoissonRegressor(), param_grid, verbose=2, cv=5)
train_and_evaluate_model(grid_poisson)
```

```
Fitting 5 folds for each of 10 candidates, totalling 50 fits
[CV] END alpha=1, fit_intercept=True, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=1, fit_intercept=True, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=1, fit_intercept=True, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=1, fit_intercept=True, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=1, fit_intercept=True, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=1, fit_intercept=False, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=1, fit_intercept=False, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=1, fit_intercept=False, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=1, fit_intercept=False, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=1, fit_intercept=False, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=True, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=True, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=True, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=True, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=True, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=1, fit_intercept=False, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=1, fit_intercept=False, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=1, fit_intercept=False, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=1, fit_intercept=False, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=1, fit_intercept=False, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=10, fit_intercept=True, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=10, fit_intercept=True, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=10, fit_intercept=True, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=10, fit_intercept=True, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=10, fit_intercept=True, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=10, fit_intercept=True, solver=lbfgs, warm_start=True; total time= 0.0s
[CV] END alpha=10, fit_intercept=True, solver=lbfgs, warm_start=True; total time= 0.0s
[CV] END alpha=10, fit_intercept=True, solver=lbfgs, warm_start=True; total time= 0.0s
[CV] END alpha=10, fit_intercept=True, solver=lbfgs, warm_start=True; total time= 0.0s
[CV] END alpha=10, fit_intercept=True, solver=lbfgs, warm_start=True; total time= 0.0s
```

```
[CV] END alpha=0.001, fit_intercept=False, solver=lbfgs, warm_start=False;
total time= 0.0s
[CV] END alpha=0.001, fit_intercept=False, solver=lbfgs, warm_start=False;
total time= 0.0s
[CV] END alpha=0.001, fit_intercept=False, solver=lbfgs, warm_start=False;
total time= 0.0s
[CV] END alpha=0.001, fit_intercept=False, solver=lbfgs, warm_start=False;
total time= 0.0s
[CV] END alpha=0.001, fit_intercept=False, solver=lbfgs, warm_start=False;
total time= 0.0s
[CV] END alpha=10, fit_intercept=True, solver=newton-cholesky, warm_start=
False; total time= 0.0s
[CV] END alpha=10, fit_intercept=True, solver=newton-cholesky, warm_start=
False; total time= 0.0s
[CV] END alpha=10, fit_intercept=True, solver=newton-cholesky, warm_start=
False; total time= 0.0s
[CV] END alpha=10, fit_intercept=True, solver=newton-cholesky, warm_start=
False; total time= 0.0s
[CV] END alpha=10, fit_intercept=True, solver=newton-cholesky, warm_start=
False; total time= 0.0s
[CV] END alpha=10, fit_intercept=True, solver=newton-cholesky, warm_start=
False; total time= 0.0s
[CV] END alpha=10, fit_intercept=False, solver=newton-cholesky, warm_start
=True; total time= 0.0s
[CV] END alpha=10, fit_intercept=False, solver=newton-cholesky, warm_start
=True; total time= 0.0s
[CV] END alpha=10, fit_intercept=False, solver=newton-cholesky, warm_start
=True; total time= 0.0s
[CV] END alpha=10, fit_intercept=False, solver=newton-cholesky, warm_start
=True; total time= 0.0s
[CV] END alpha=10, fit_intercept=False, solver=newton-cholesky, warm_start
=True; total time= 0.0s
[CV] END alpha=10, fit_intercept=False, solver=newton-cholesky, warm_start
=True; total time= 0.0s
[CV] END alpha=1, fit_intercept=True, solver=lbfgs, warm_start=True; total
time= 0.0s
[CV] END alpha=1, fit_intercept=True, solver=lbfgs, warm_start=True; total
time= 0.0s
[CV] END alpha=1, fit_intercept=True, solver=lbfgs, warm_start=True; total
time= 0.0s
[CV] END alpha=1, fit_intercept=True, solver=lbfgs, warm_start=True; total
time= 0.0s
[CV] END alpha=1, fit_intercept=True, solver=lbfgs, warm_start=True; total
time= 0.0s
Mean Absolute Percentage Error: 0.06294123795511576
Root Mean Squared Error: 135479.5624858958
Root Mean Squared Log Error: 0.08324169711154629
R2 Score: 0.9553134871395782
```

In [149]:

```
grid_poisson.best_params_
```

Out[149]:

```
{'warm_start': True,
 'solver': 'newton-cholesky',
 'fit_intercept': True,
 'alpha': 10}
```

In [150]:

```
param_grid = {'solver': ['lbfgs', 'newton-cholesky'],
              'fit_intercept': [True, False],
              'alpha': [0.001, 0.01, 0.1, 1, 10],
              'warm_start': [True, False]}

grid_gamma = RandomizedSearchCV(GammaRegressor(), param_grid, verbose=2, cv=5)
train_and_evaluate_model(grid_gamma)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[CV] END alpha=0.001, fit_intercept=True, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=True, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=True, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=True, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=True, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=1, fit_intercept=True, solver=lbgfs, warm_start=True; total time= 0.0s
[CV] END alpha=1, fit_intercept=True, solver=lbgfs, warm_start=True; total time= 0.0s
[CV] END alpha=1, fit_intercept=True, solver=lbgfs, warm_start=True; total time= 0.0s
[CV] END alpha=1, fit_intercept=True, solver=lbgfs, warm_start=True; total time= 0.0s
[CV] END alpha=1, fit_intercept=True, solver=lbgfs, warm_start=True; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=True, solver=lbgfs, warm_start=True; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=True, solver=lbgfs, warm_start=True; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=True, solver=lbgfs, warm_start=True; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=True, solver=lbgfs, warm_start=True; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=True, solver=lbgfs, warm_start=True; total time= 0.0s
[CV] END alpha=0.1, fit_intercept=True, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=0.1, fit_intercept=True, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=0.1, fit_intercept=True, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=0.1, fit_intercept=True, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=0.1, fit_intercept=True, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=0.1, fit_intercept=True, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=False, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=False, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=False, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=False, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=False, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=10, fit_intercept=True, solver=newton-cholesky, warm_start=False; total time= 0.0s
[CV] END alpha=10, fit_intercept=True, solver=newton-cholesky, warm_start=False; total time= 0.0s
[CV] END alpha=10, fit_intercept=True, solver=newton-cholesky, warm_start=False; total time= 0.0s
[CV] END alpha=10, fit_intercept=True, solver=newton-cholesky, warm_start=False; total time= 0.0s
[CV] END alpha=10, fit_intercept=True, solver=newton-cholesky, warm_start=False; total time= 0.0s
```

```
[CV] END alpha=0.1, fit_intercept=True, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=0.1, fit_intercept=True, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=0.1, fit_intercept=True, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=0.1, fit_intercept=True, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=0.1, fit_intercept=True, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=0.01, fit_intercept=True, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=0.01, fit_intercept=True, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=0.01, fit_intercept=True, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=0.01, fit_intercept=True, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=0.01, fit_intercept=True, solver=newton-cholesky, warm_start=True; total time= 0.0s
[CV] END alpha=0.01, fit_intercept=True, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=0.01, fit_intercept=True, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=0.01, fit_intercept=True, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=0.01, fit_intercept=True, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=0.01, fit_intercept=True, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=False, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=False, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=False, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=False, solver=lbfgs, warm_start=False; total time= 0.0s
[CV] END alpha=0.001, fit_intercept=False, solver=lbfgs, warm_start=False; total time= 0.0s
Mean Absolute Percentage Error: 0.05874172503527488
Root Mean Squared Error: 135231.56435136718
Root Mean Squared Log Error: 0.07756227647474297
R2 Score: 0.9554769365713998
```

In [151]:

```
grid_gamma.best_params_
```

Out[151]:

```
{'warm_start': True, 'solver': 'lbfgs', 'fit_intercept': True, 'alpha': 0.001}
```

In [152]:

```
param_grid = {'alpha_1': [1e-5,1e-6,1e-7,1e-8],
              'alpha_2': [1e-5,1e-6,1e-7,1e-8],
              'lambda_1': [1e-5,1e-6,1e-7,1e-8],
              'lambda_2': [1e-5,1e-6,1e-7,1e-8],
              'fit_intercept': [True,False],
              'compute_score': [True,False]}

grid_ard = RandomizedSearchCV(ARDRegression(),param_grid,verbose=2,cv=5)
train_and_evaluate_model(grid_ard)
```


97/105

```
[CV] END alpha_1=1e-08, alpha_2=1e-05, compute_score=True, fit_intercept=F
alse, lambda_1=1e-07, lambda_2=1e-06; total time= 0.0s
[CV] END alpha_1=1e-08, alpha_2=1e-05, compute_score=True, fit_intercept=F
alse, lambda_1=1e-07, lambda_2=1e-06; total time= 0.0s
[CV] END alpha_1=1e-08, alpha_2=1e-05, compute_score=True, fit_intercept=F
alse, lambda_1=1e-07, lambda_2=1e-06; total time= 0.0s
[CV] END alpha_1=1e-08, alpha_2=1e-05, compute_score=True, fit_intercept=F
alse, lambda_1=1e-07, lambda_2=1e-06; total time= 0.0s
[CV] END alpha_1=1e-08, alpha_2=1e-05, compute_score=True, fit_intercept=F
alse, lambda_1=1e-07, lambda_2=1e-06; total time= 0.0s
[CV] END alpha_1=1e-07, alpha_2=1e-05, compute_score=True, fit_intercept=T
rue, lambda_1=1e-05, lambda_2=1e-06; total time= 0.0s
[CV] END alpha_1=1e-07, alpha_2=1e-05, compute_score=True, fit_intercept=T
rue, lambda_1=1e-05, lambda_2=1e-06; total time= 0.0s
[CV] END alpha_1=1e-07, alpha_2=1e-05, compute_score=True, fit_intercept=T
rue, lambda_1=1e-05, lambda_2=1e-06; total time= 0.0s
[CV] END alpha_1=1e-07, alpha_2=1e-05, compute_score=True, fit_intercept=T
rue, lambda_1=1e-05, lambda_2=1e-06; total time= 0.0s
[CV] END alpha_1=1e-07, alpha_2=1e-05, compute_score=True, fit_intercept=T
rue, lambda_1=1e-05, lambda_2=1e-06; total time= 0.0s
[CV] END alpha_1=1e-08, alpha_2=1e-06, compute_score=True, fit_intercept=F
alse, lambda_1=1e-07, lambda_2=1e-06; total time= 0.0s
[CV] END alpha_1=1e-08, alpha_2=1e-06, compute_score=True, fit_intercept=F
alse, lambda_1=1e-07, lambda_2=1e-06; total time= 0.0s
[CV] END alpha_1=1e-08, alpha_2=1e-06, compute_score=True, fit_intercept=F
alse, lambda_1=1e-07, lambda_2=1e-06; total time= 0.0s
[CV] END alpha_1=1e-08, alpha_2=1e-06, compute_score=True, fit_intercept=F
alse, lambda_1=1e-07, lambda_2=1e-06; total time= 0.0s
[CV] END alpha_1=1e-08, alpha_2=1e-06, compute_score=True, fit_intercept=F
alse, lambda_1=1e-07, lambda_2=1e-06; total time= 0.0s
[CV] END alpha_1=1e-08, alpha_2=1e-07, compute_score=False, fit_intercept=
True, lambda_1=1e-06, lambda_2=1e-07; total time= 0.0s
[CV] END alpha_1=1e-08, alpha_2=1e-07, compute_score=False, fit_intercept=
True, lambda_1=1e-06, lambda_2=1e-07; total time= 0.0s
[CV] END alpha_1=1e-08, alpha_2=1e-07, compute_score=False, fit_intercept=
True, lambda_1=1e-06, lambda_2=1e-07; total time= 0.0s
[CV] END alpha_1=1e-08, alpha_2=1e-07, compute_score=False, fit_intercept=
True, lambda_1=1e-06, lambda_2=1e-07; total time= 0.0s
[CV] END alpha_1=1e-08, alpha_2=1e-07, compute_score=False, fit_intercept=
True, lambda_1=1e-06, lambda_2=1e-07; total time= 0.0s
Mean Absolute Percentage Error: 0.03767929032989868
Root Mean Squared Error: 92664.15716479663
Root Mean Squared Log Error: 0.054885830451669886
R2 Score: 0.9790948901271462
```

In [153]:

```
grid_ard.best_params_
```

Out[153]:

```
{'lambda_2': 1e-06,
 'lambda_1': 1e-08,
 'fit_intercept': True,
 'compute_score': False,
 'alpha_2': 1e-07,
 'alpha_1': 1e-08}
```

In [154]:

```
param_grid = {'fit_intercept': [True, False],
              'positive': [True, False]}

grid_lr = RandomizedSearchCV(LinearRegression(), param_grid, verbose=2, cv=5)
train_and_evaluate_model(grid_lr)
```

Fitting 5 folds for each of 4 candidates, totalling 20 fits

```
[CV] END .....fit_intercept=True, positive=True; total time=
0.0s
[CV] END .....fit_intercept=True, positive=True; total time=
0.0s
[CV] END .....fit_intercept=True, positive=True; total time=
0.0s
[CV] END .....fit_intercept=True, positive=True; total time=
0.0s
[CV] END .....fit_intercept=True, positive=True; total time=
0.0s
[CV] END .....fit_intercept=True, positive=False; total time=
0.0s
[CV] END .....fit_intercept=True, positive=False; total time=
0.0s
[CV] END .....fit_intercept=True, positive=False; total time=
0.0s
[CV] END .....fit_intercept=True, positive=False; total time=
0.0s
[CV] END .....fit_intercept=True, positive=False; total time=
0.0s
[CV] END .....fit_intercept=False, positive=True; total time=
0.0s
[CV] END .....fit_intercept=False, positive=True; total time=
0.0s
[CV] END .....fit_intercept=False, positive=True; total time=
0.0s
[CV] END .....fit_intercept=False, positive=True; total time=
0.0s
[CV] END .....fit_intercept=False, positive=True; total time=
0.0s
[CV] END .....fit_intercept=False, positive=False; total time=
0.0s
[CV] END .....fit_intercept=False, positive=False; total time=
0.0s
[CV] END .....fit_intercept=False, positive=False; total time=
0.0s
[CV] END .....fit_intercept=False, positive=False; total time=
0.0s
```

Mean Absolute Percentage Error: 0.037602824334626514

Root Mean Squared Error: 95719.65932929565

Root Mean Squared Log Error: 0.054083587741203705

R2 Score: 0.9776935128173896

In [155]:

```
grid_lr.best_params_
```

Out[155]:

```
{'positive': False, 'fit_intercept': True}
```

Hyperparameter-Optimized Model Performance Comparison

In [156]:

```
model_perfs = pd.DataFrame({'Model': models, 'MAPE': mape_scores, 'RMSE': rmse_scores, '  
model_perfs = model_perfs.sort_values('R2',ascending=False)  
model_perfs
```

Out[156]:

	Model	MAPE	RMSE	RMSLE	
28	VotingRegressor(estimators=[('CAT',ln ...	0.024105	7.483545e+04	0.037744	0
39	RandomizedSearchCV(estimator=ExtraTreesRegress...	0.024178	7.749662e+04	0.039357	0
21	([DecisionTreeRegressor(criterion='friedman_ms...	0.025872	7.938804e+04	0.040228	0
20	(ExtraTreeRegressor(random_state=2090432581), ...	0.020761	8.088466e+04	0.041309	0
47	RandomizedSearchCV(cv=5, estimator=ElasticNetC...	0.037549	8.923582e+04	0.053935	0
45	RandomizedSearchCV(cv=5, estimator=LassoCV(),\...	0.036074	8.948055e+04	0.052035	0
51	RandomizedSearchCV(cv=5, estimator=ARDRegressi...	0.037679	9.266416e+04	0.054886	0
19	(DecisionTreeRegressor(random_state=174724458)...	0.032294	9.331461e+04	0.045934	0
18	(DecisionTreeRegressor(max_features=1.0, rando...	0.031221	9.415250e+04	0.044139	0
5	SGDRegressor()	0.037567	9.439265e+04	0.053727	0
2	LassoCV()	0.037831	9.461667e+04	0.054418	0
37	RandomizedSearchCV(cv=5, estimator=BaggingRegr...	0.031812	9.470061e+04	0.045514	0
1	RidgeCV()	0.037577	9.545932e+04	0.053988	0
46	RandomizedSearchCV(cv=5, estimator=RidgeCV(),\...	0.037577	9.545932e+04	0.053988	0
11	RANSACRegressor()	0.037603	9.571966e+04	0.054084	0
0	LinearRegression()	0.037603	9.571966e+04	0.054084	0
52	RandomizedSearchCV(cv=5, estimator=LinearRegre...	0.037603	9.571966e+04	0.054084	0
40	RandomizedSearchCV(cv=5,\n e...	0.037618	9.574424e+04	0.054104	0
41	RandomizedSearchCV(cv=5, estimator=SGDRegresso...	0.037189	9.587206e+04	0.053823	0
29	RandomizedSearchCV(cv=5, estimator=RandomFores...	0.032077	9.628709e+04	0.047403	0
6	HuberRegressor()	0.032574	9.771381e+04	0.054128	0
12	TheilSenRegressor()	0.030573	1.001043e+05	0.051125	0
33	RandomizedSearchCV(cv=RepeatedKfold(n_repeats=...	0.043065	1.029213e+05	0.062062	0
44	RandomizedSearchCV(cv=5, estimator=PassiveAggr...	0.034113	1.030828e+05	0.051854	0
26	<catboost.core.CatBoostRegressor object at 0x0...	0.035152	1.087557e+05	0.061378	0
35	RandomizedSearchCV(cv=6, estimator=GradientBoo...	0.049553	1.244436e+05	0.064934	0
48	RandomizedSearchCV(cv=5, estimator=AdaBoostReg...	0.063824	1.253858e+05	0.075975	0
27	LGBMRegressor()	0.044858	1.262488e+05	0.064916	0
34	RandomizedSearchCV(cv=RepeatedKfold(n_repeats=...	0.044681	1.281810e+05	0.067728	0
25	XGBRegressor(base_score=None, booster=None, ca...	0.042063	1.283637e+05	0.060352	0
22	(DecisionTreeRegressor(max_depth=3, random_sta...	0.060388	1.307803e+05	0.076905	0
23	HistGradientBoostingRegressor()	0.048545	1.338957e+05	0.067489	0
36	RandomizedSearchCV(cv=5,\n e...	0.044126	1.349660e+05	0.069766	0
50	RandomizedSearchCV(cv=5, estimator=GammaRegres...	0.058742	1.352316e+05	0.077562	0
49	RandomizedSearchCV(cv=5, estimator=PoissonRegr...	0.062941	1.354796e+05	0.083242	0
8	PoissonRegressor()	0.062941	1.354802e+05	0.083241	0
43	RandomizedSearchCV(cv=5, estimator=TweedieRegr...	0.059141	1.360808e+05	0.077778	0

	Model	MAPE	RMSE	RMSLE	
31	RandomizedSearchCV(cv=5, estimator=HistGradien...	0.050687	1.376967e+05	0.078475	0
38	RandomizedSearchCV(estimator=LGBMRegressor()),\...	0.052871	1.422376e+05	0.074401	0
30	RandomizedSearchCV(cv=RepeatedKFold(n_repeats=...	0.081620	1.754687e+05	0.105170	0
17	DecisionTreeRegressor()	0.057474	1.780531e+05	0.089337	0
14	KNeighborsRegressor()	0.096216	1.942969e+05	0.118579	0
10	TweedieRegressor()	0.122175	2.034528e+05	0.153165	0
7	GammaRegressor()	0.128221	2.267689e+05	0.157562	0
3	ElasticNetCV()	0.387979	6.378490e+05	0.398606	0
32	RandomizedSearchCV(cv=RepeatedKFold(n_repeats=...	0.392319	6.393462e+05	0.401252	0
42	RandomizedSearchCV(cv=5, estimator=HuberRegres...	0.388217	6.408426e+05	0.399389	0
13	ARDRegression()	0.389750	6.408954e+05	0.400184	-0
15	SVR()	0.393779	6.411845e+05	0.402406	-0
9	QuantileRegressor()	0.393789	6.411975e+05	0.402414	-0
4	PassiveAggressiveRegressor()	0.907267	1.796564e+06	2.491688	-6
16	LinearSVR()	0.999907	1.932580e+06	9.369523	-8
24	MLPRegressor()	0.999992	1.932699e+06	11.735068	-8

After completion of the entire hyperparameter tuning process, it is clearly evident that the Voting Regressor outperforms all the other models by obtaining an astounding r2 score of approximately 98.64%.



Saving the best performing model for future use

In [158]:

```
vc = VotingRegressor(estimators=[
    ('CAT', CatBoostRegressor(silent=True)),
    ('LR', LinearRegression()),
    ('SGD', SGDRegressor()),
    ('GB', GradientBoostingRegressor()),
    ('ET', ExtraTreesRegressor()),
    ('BAG', BaggingRegressor()),
    ('RF', RandomForestRegressor()),
    ('THEIL', TheilSenRegressor()),
    ('RANSAC', RANSACRegressor()),
    ('HUBER', HuberRegressor()),
    ('LASSO', LassoCV()),
    ('RIDGE', RidgeCV())
])
vc.fit(X_train_final, y_train)
```

Out[158]:

```
VotingRegressor(estimators=[('CAT',
                             <catboost.core.CatBoostRegressor object at 0x
000002D85DE76530>),
                             ('LR', LinearRegression()), ('SGD', SGDRegress
or()),
                             ('GB', GradientBoostingRegressor()),
                             ('ET', ExtraTreesRegressor()),
                             ('BAG', BaggingRegressor()),
                             ('RF', RandomForestRegressor()),
                             ('THEIL', TheilSenRegressor()),
                             ('RANSAC', RANSACRegressor()),
                             ('HUBER', HuberRegressor()), ('LASSO', LassoCV
()),
                             ('RIDGE', RidgeCV())])
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [159]:

```
vc_pred = vc.predict(X_test_final)
print("R2:", r2_score(y_test, vc_pred))
```

R2: 0.9851030297616545

In [160]:

```
joblib.dump(vc, 'sales_predictor.h5')
```

Out[160]:

```
['sales_predictor.h5']
```


In [161]:

```
model = joblib.load('sales_predictor.h5')
model
```

Out[161]:

```
VotingRegressor(estimators=[('CAT',
                             <catboost.core.CatBoostRegressor object at 0x
000002D85E465E70>),
                             ('LR', LinearRegression()), ('SGD', SGDRegress
or()),
                             ('GB', GradientBoostingRegressor()),
                             ('ET', ExtraTreesRegressor()),
                             ('BAG', BaggingRegressor()),
                             ('RF', RandomForestRegressor()),
                             ('THEIL', TheilSenRegressor()),
                             ('RANSAC', RANSACRegressor()),
                             ('HUBER', HuberRegressor()), ('LASSO', LassoCV
()),
                             ('RIDGE', RidgeCV())])
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.